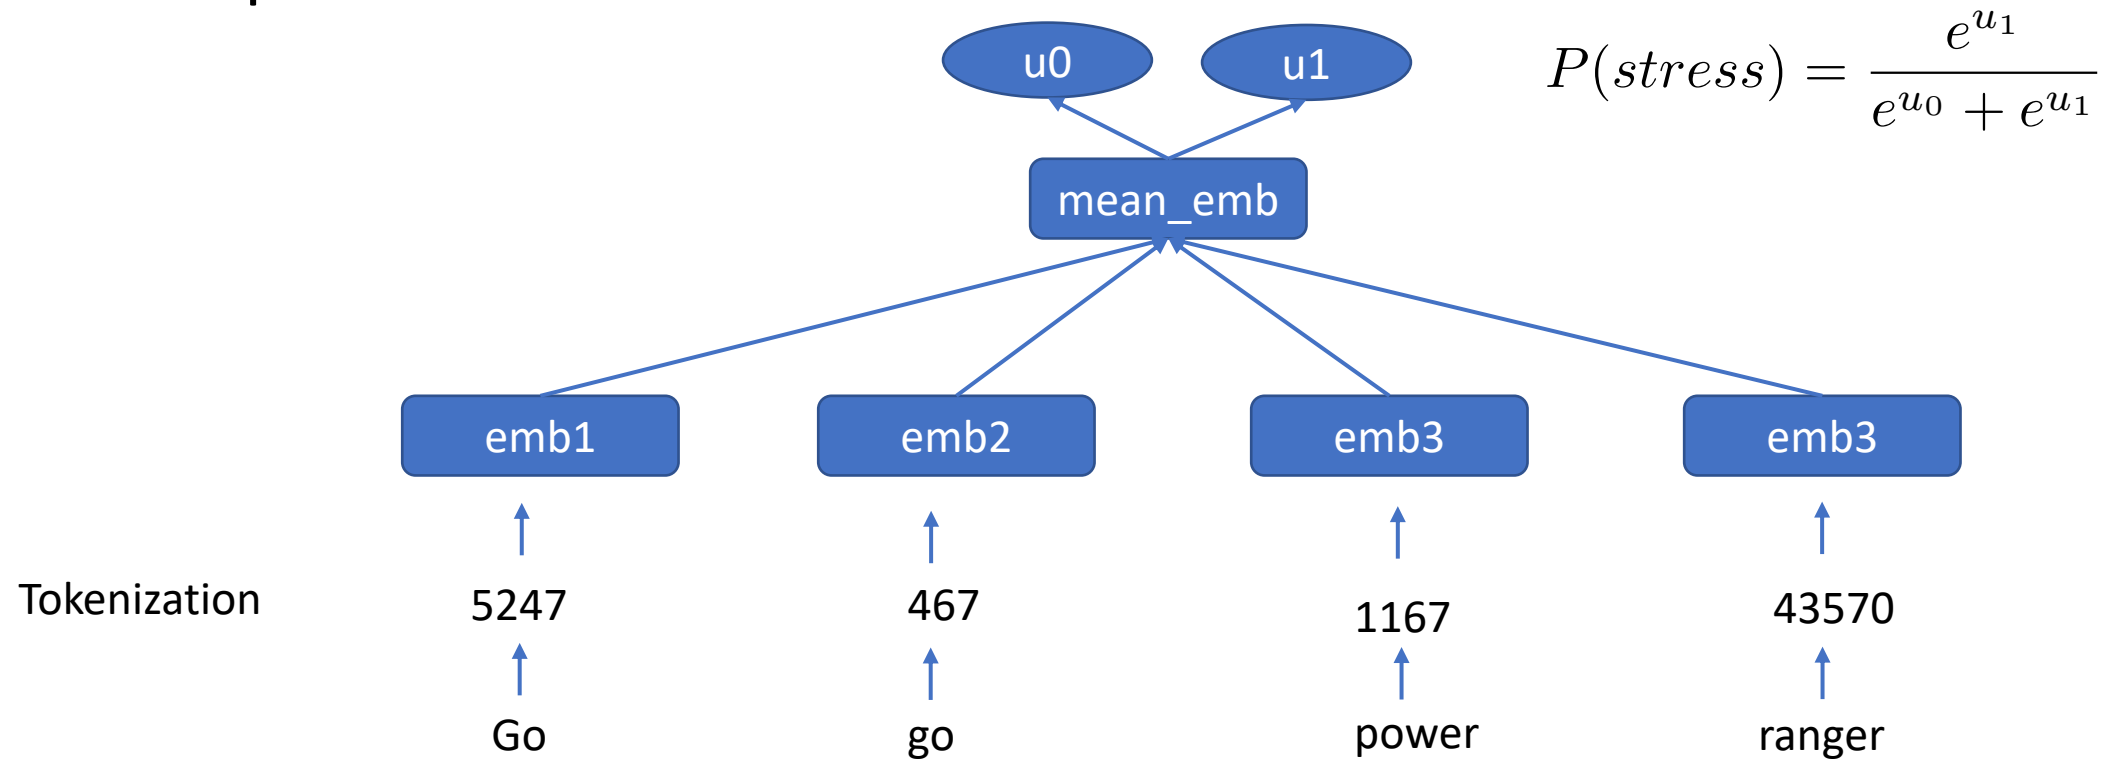


# Gen-AI: Technical and Social

## Lecture 03: Transformer and Text Generation

# Recap: Text Classification

- A simple model



# Recap: Text Classification

- A simple model

```
class EmbeddingModel(nn.Module):
    def __init__(self, vocab_size, embedding_dim, hidden_dim, num_classes):
        super().__init__()
        self.embedding = nn.Embedding(vocab_size, embedding_dim)
        self.fc1 = nn.Linear(embedding_dim, num_classes)

    def forward(self, inputs):
        output = self.embedding(inputs)
        output = torch.mean(output, dim=1) # Mean pooling
        output = self.fc1(output)
        return output

model = EmbeddingModel(tokenizer.vocab_size, 128, 128, 2)
```

# Issues of Word Embeddings?

- 1. Does not take the surrounding contexts into account
- 2. Does not take care of the order of the words

# Context matters

Word	Example Contexts
it	The animal didn't cross the street because <b>it</b> was too tired The animal didn't cross the street because <b>it</b> was too wide
station	The train left the <b>station</b> on time The radio <b>station</b> was playing 60s hits I was <b>stationed</b> on a remote island in Polynesia

# Order matter

She only eats pizza -> Only she eats pizza

The dog chased the cat -> The cat chased the dog.

....

# Transformer

By failing to prepare, you are preparing to fail.

---Benjamin Franklin

Let's prepare a right architecture...

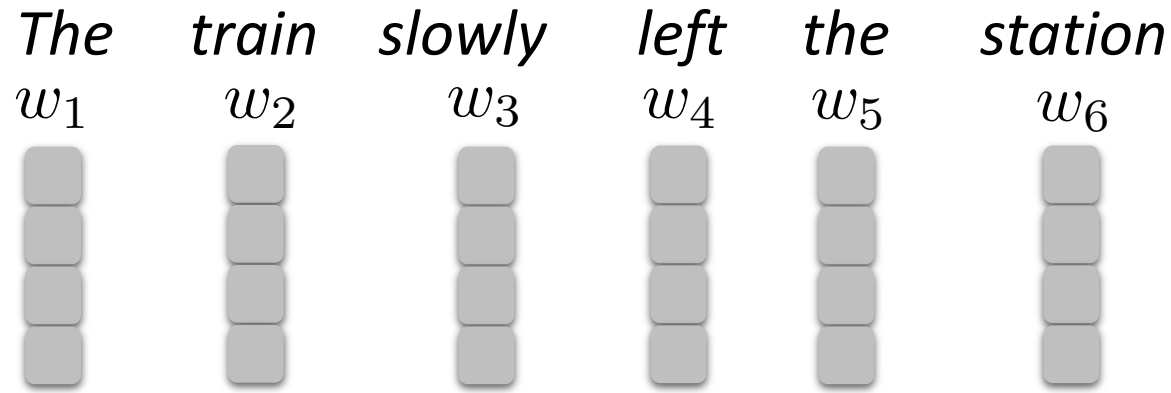
We will focus on this first

How to take the surrounding context of  
each word into account

How to take the order of the words into  
account

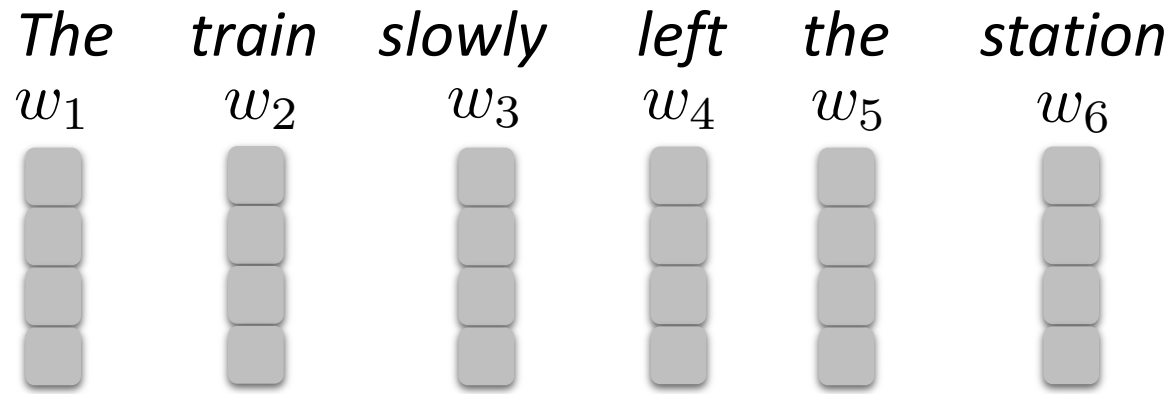


# From embeddings to *contextual embeddings*



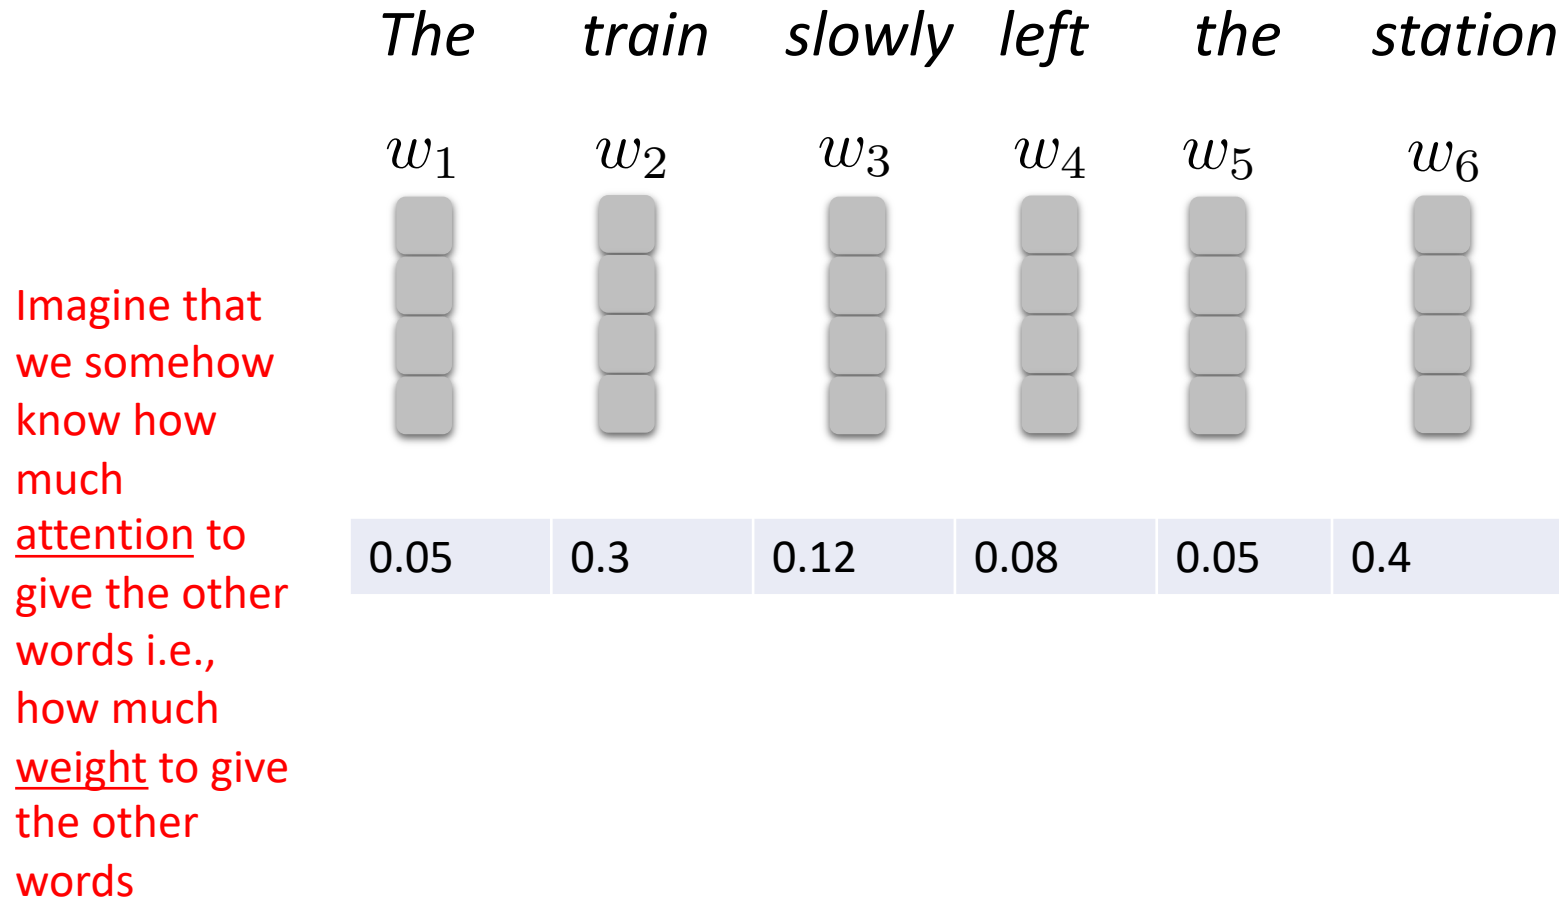
- We can easily get **stand-alone embeddings** for all the words

# From embeddings to *contextual embeddings*

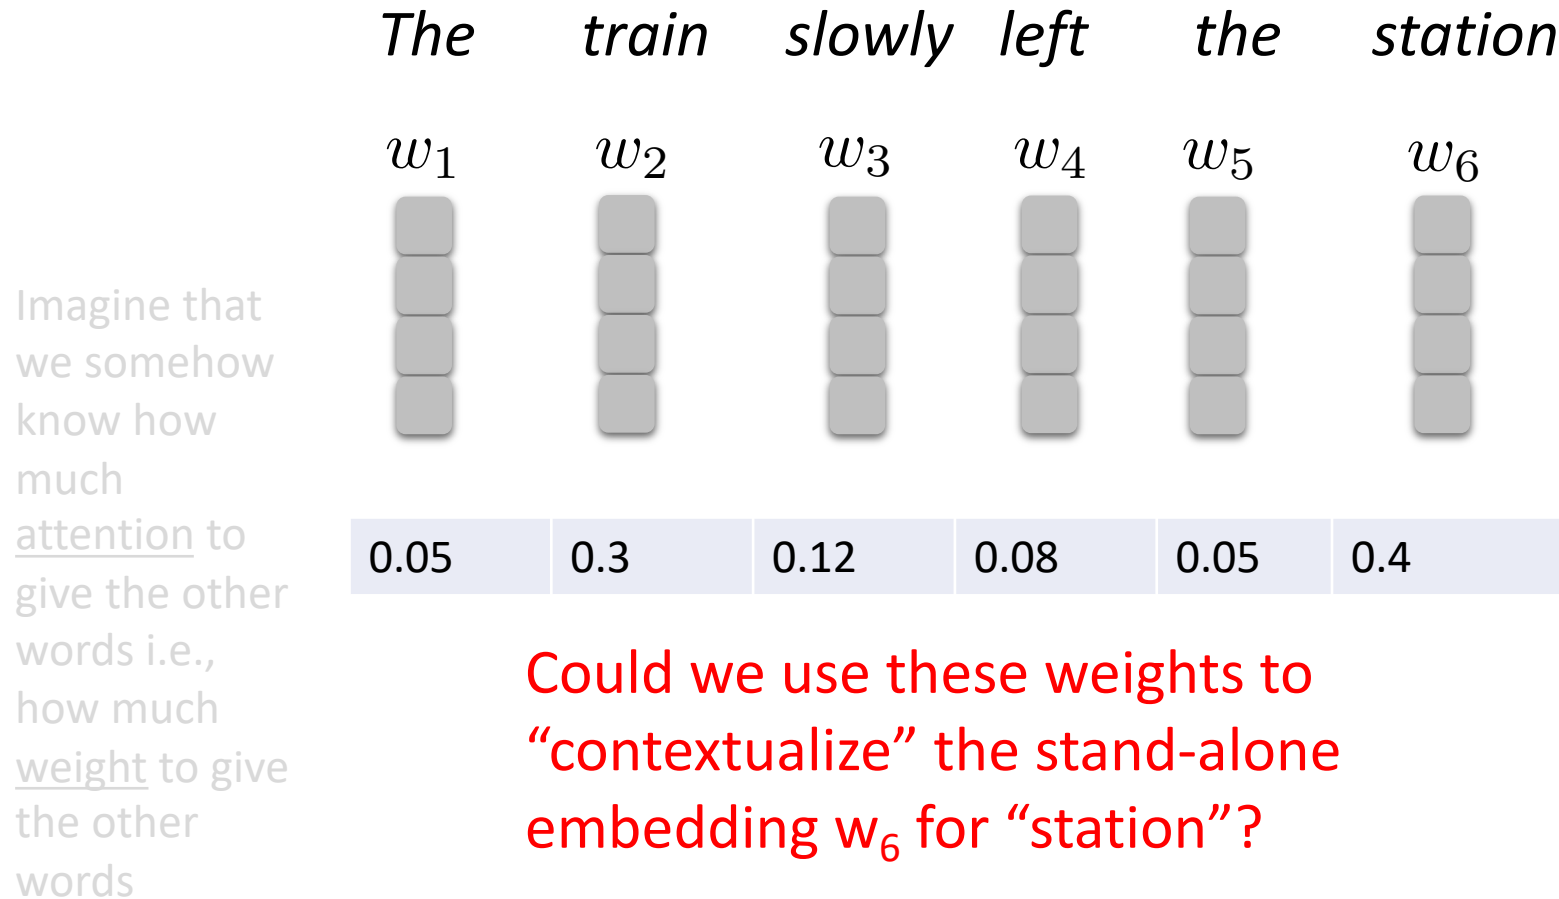


- We can easily get stand-alone embeddings for all the words
- How can we modify station's embedding so that it incorporates the other words?

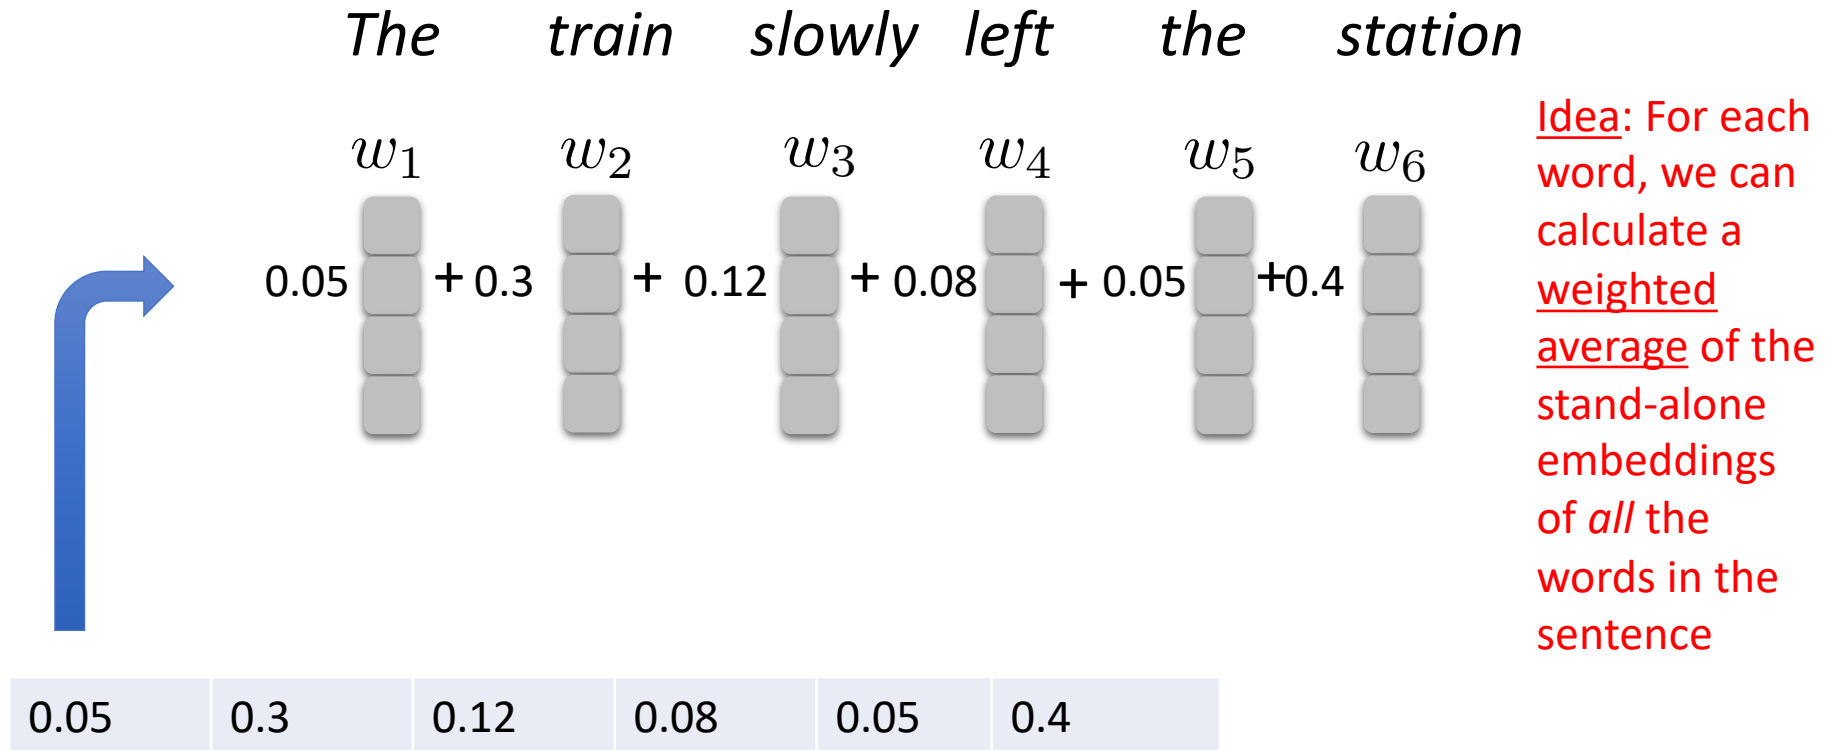
# From embeddings to *contextual embeddings*



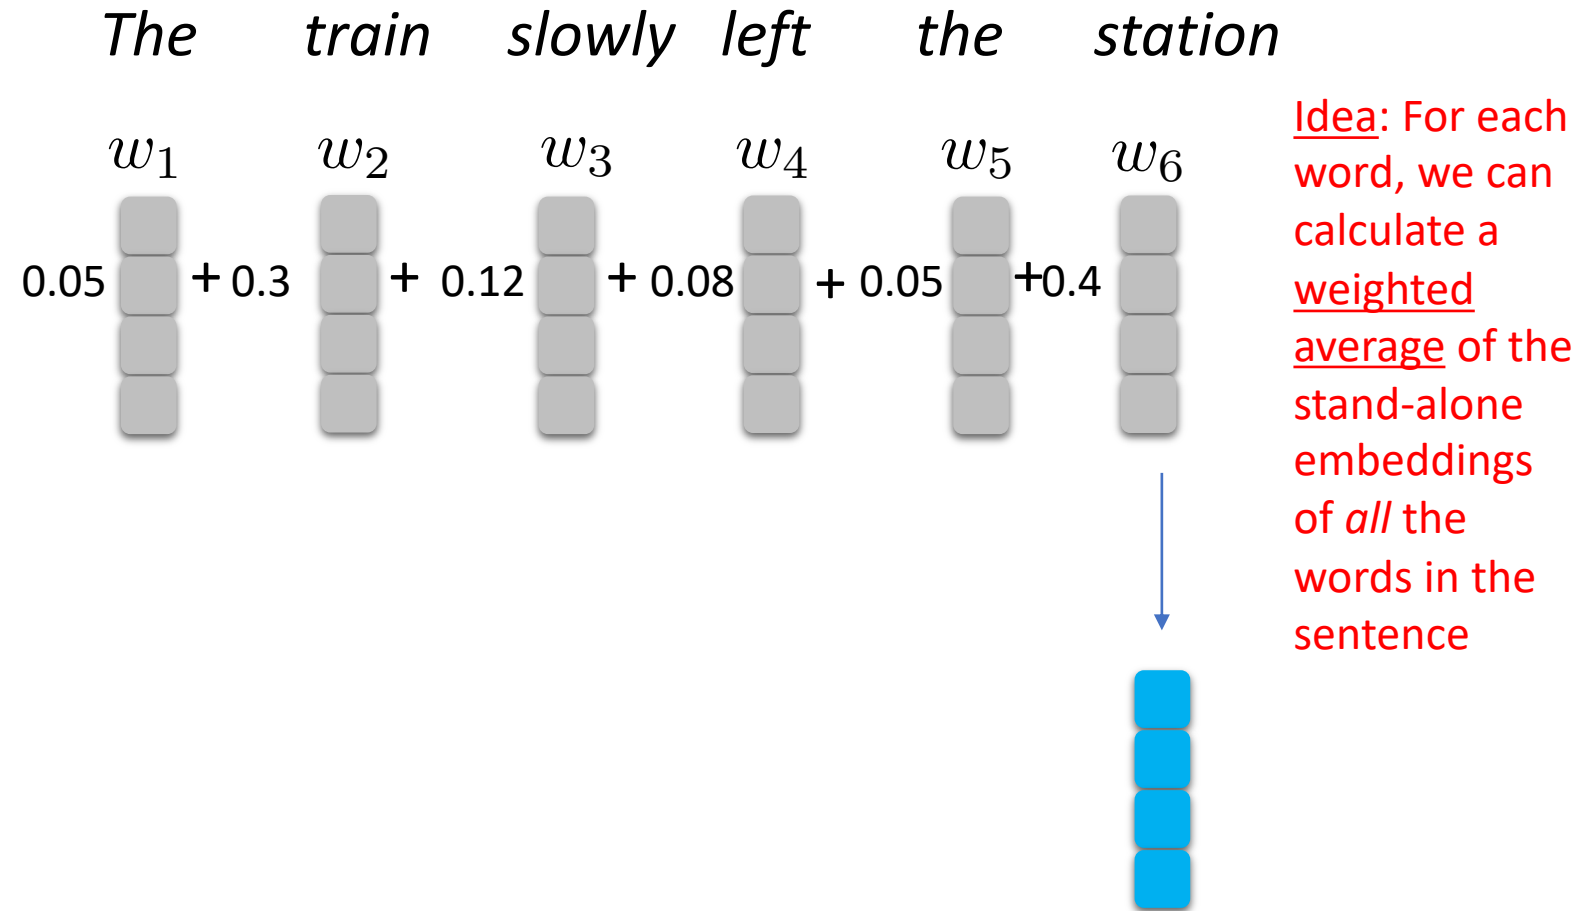
# From embeddings to *contextual embeddings*



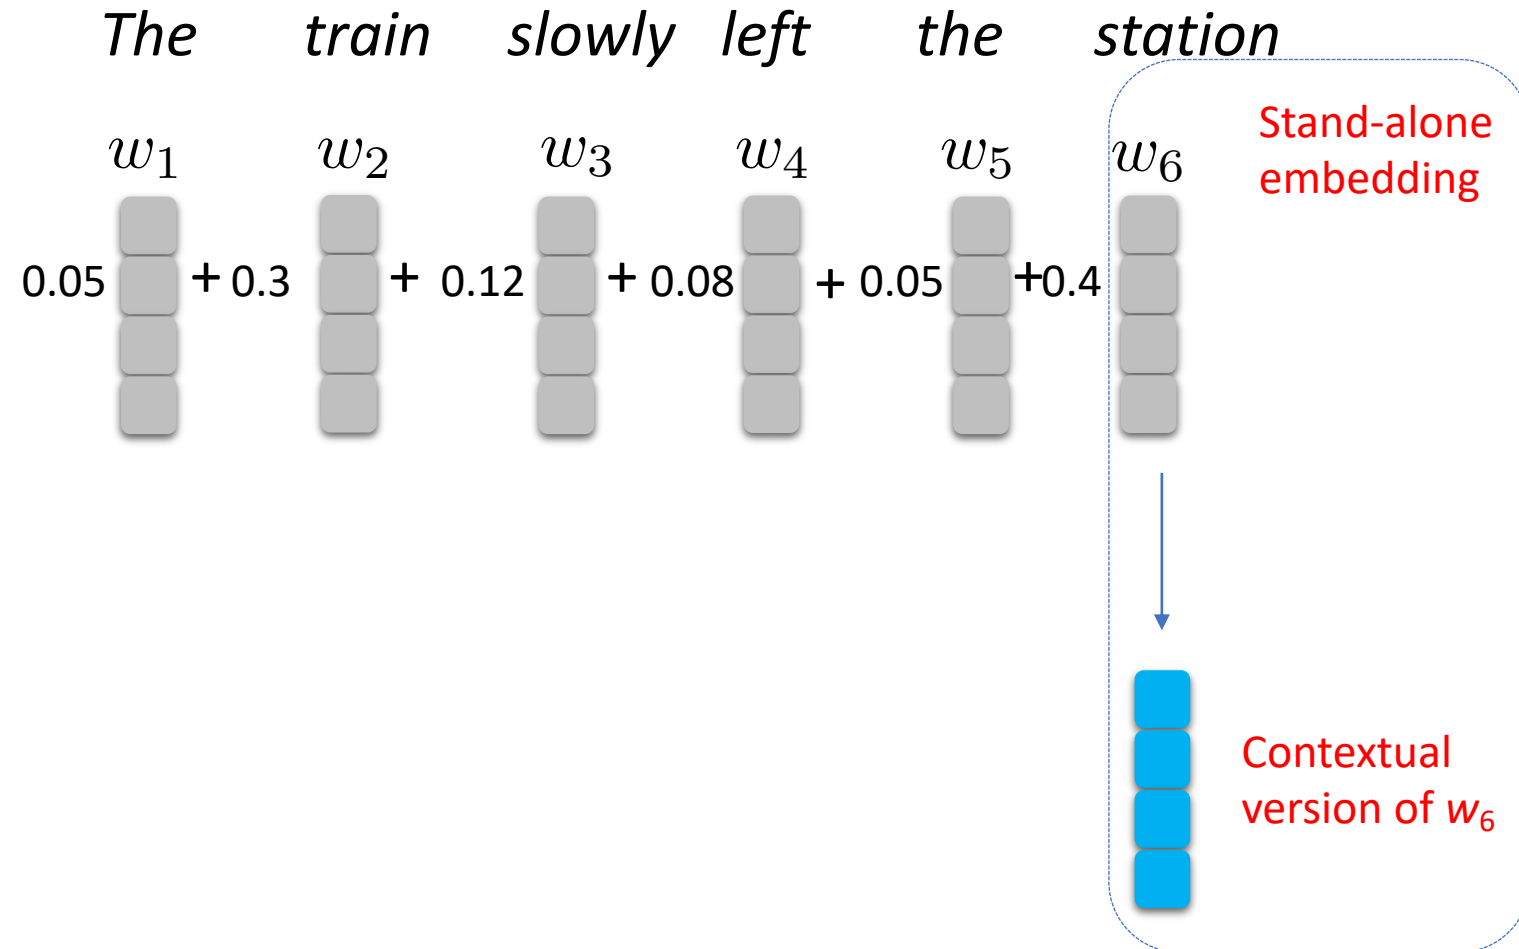
# From embeddings to *contextual embeddings*



# From embeddings to *contextual embeddings*

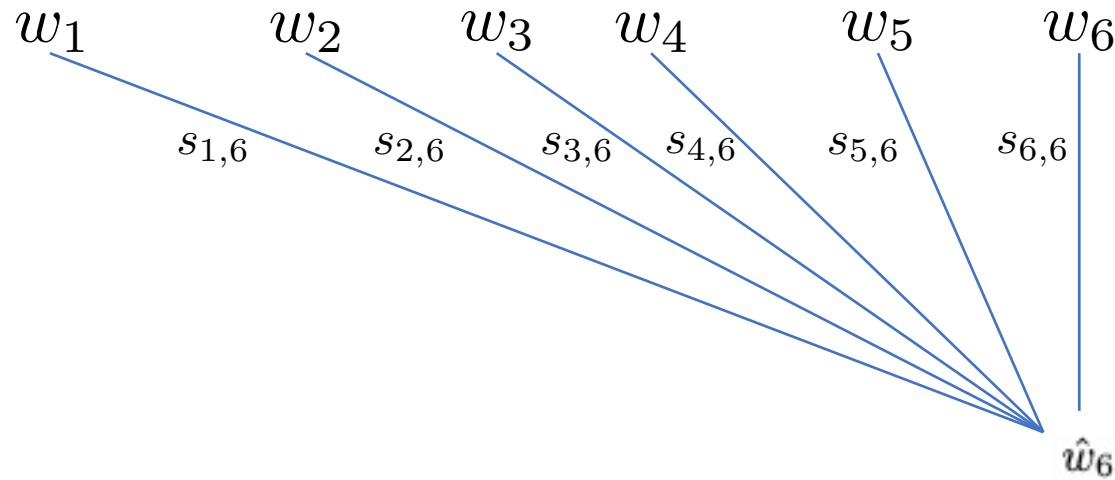


# From embeddings to *contextual embeddings*



# Same thing but more abstract

*The train slowly left the station*



Stand-alone  
embedding

$$\hat{w}_6 = s_{1,6}w_1 + s_{2,6}w_2 + s_{3,6}w_3 + s_{4,6}w_4 + s_{5,6}w_5 + s_{6,6}w_6$$

Contextual  
version of  $w_6$



For a given word (e.g., 'station'), how should the weights of the other words be chosen?

For a given word (e.g., ‘station’), how should the weights of the other words be chosen?

## Intuition

- The weight of a word should be proportional to how related it is to the word “station”

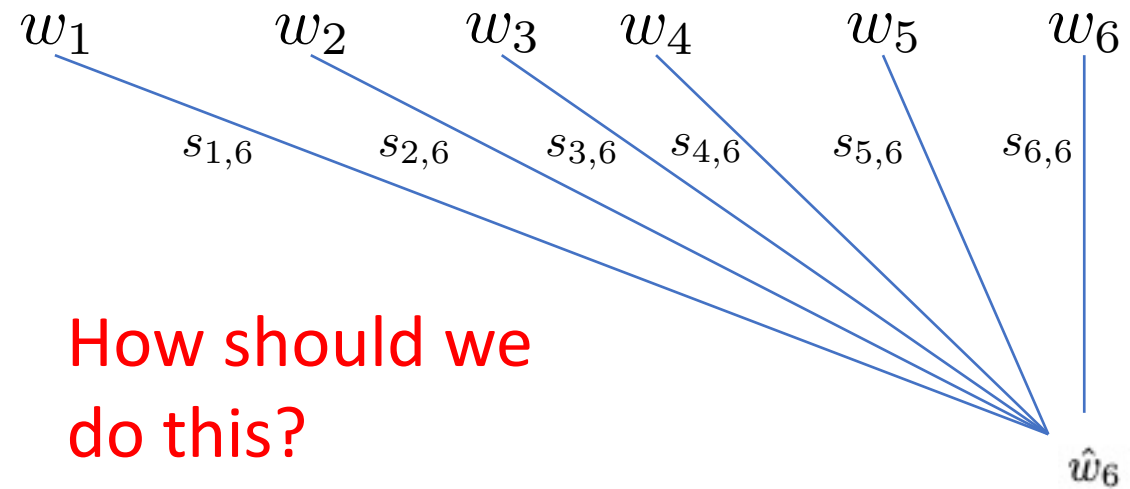
# For a given word (e.g., ‘station’), how should the weights of the other words be chosen?

## Intuition

- The weight of a word should be proportional to how related it is to the word “station”
- One way to quantify how “related” two words are: the *dot-product* of their stand-alone embeddings

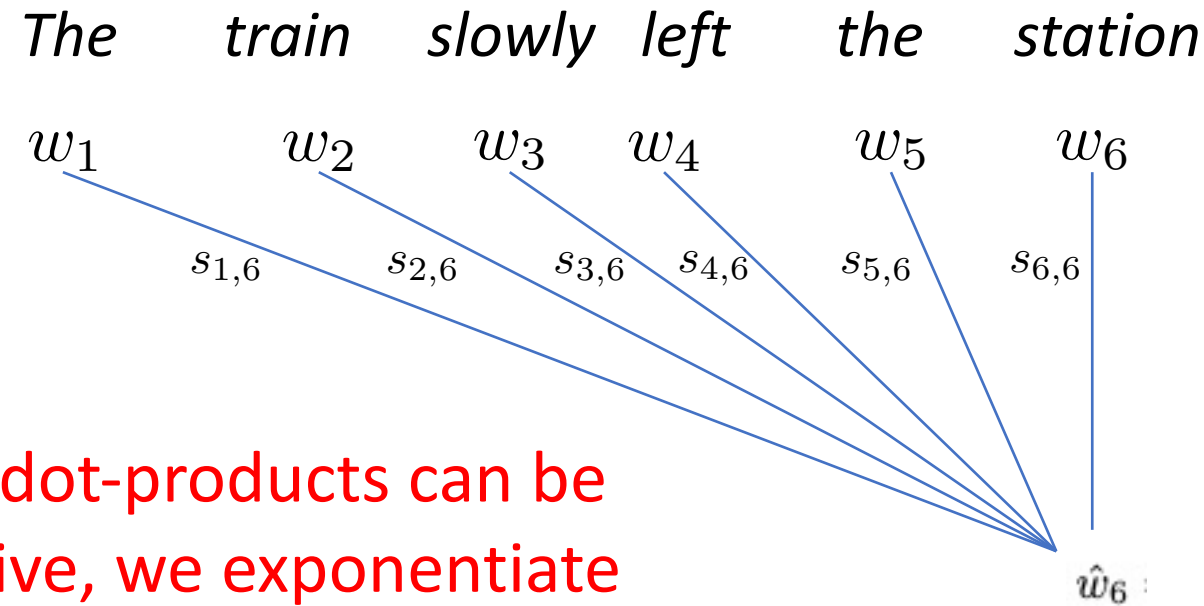
We need to convert dot-products to proper weights\*

*The train slowly left the station*



\* non-negative, and summing to 1.0

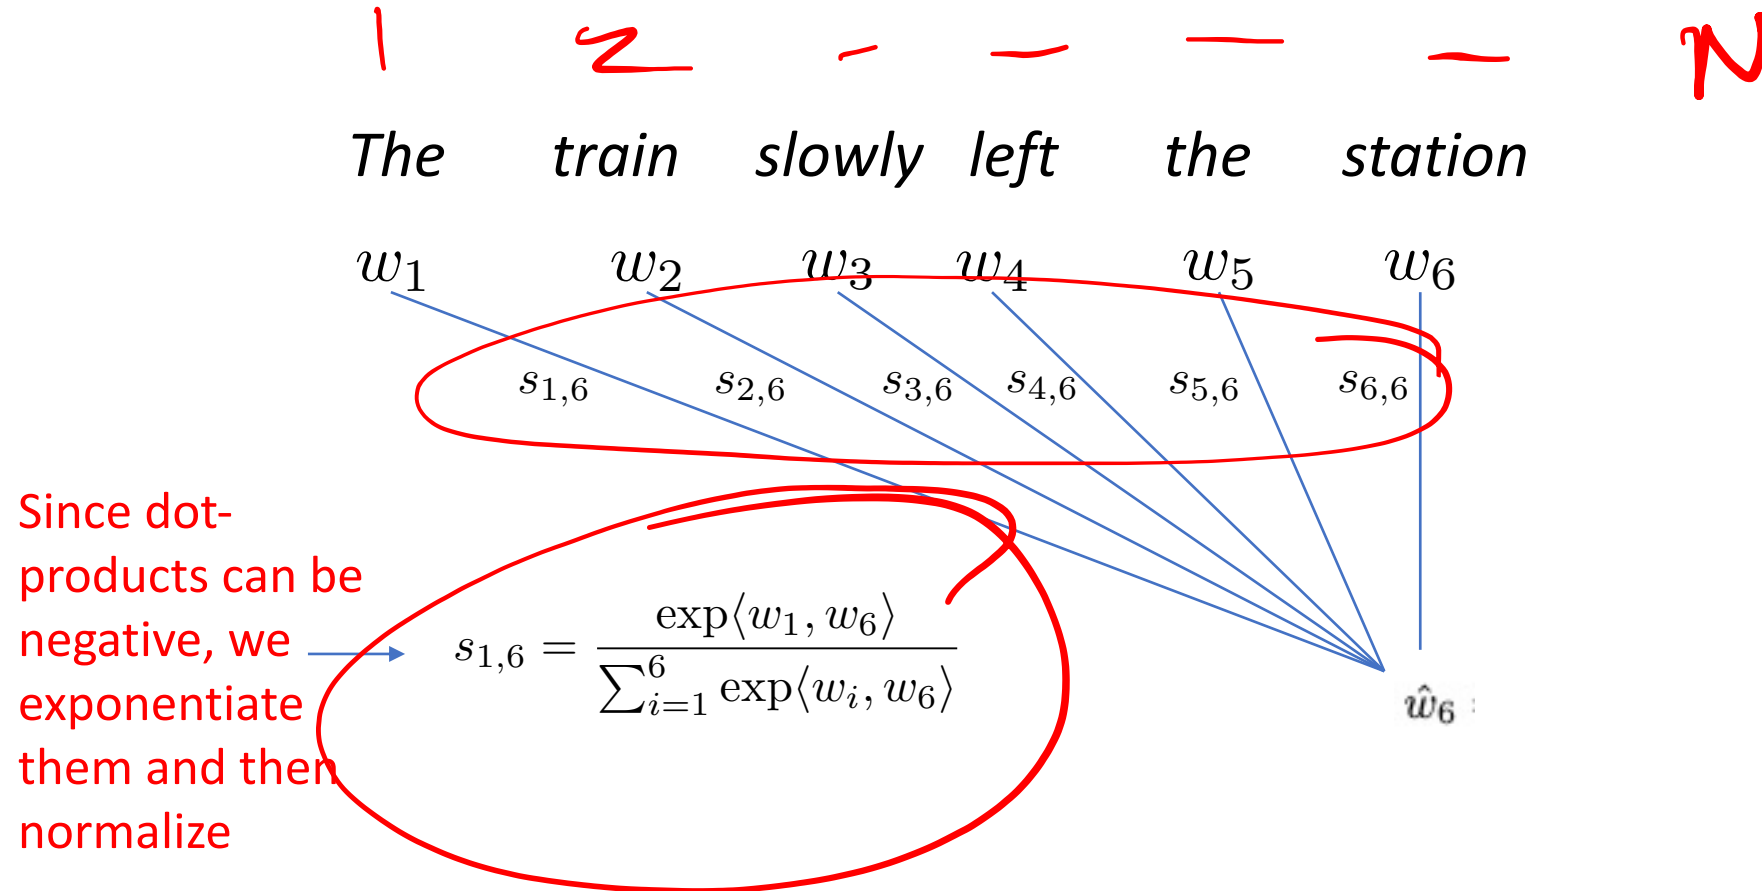
We need to convert dot-products to proper weights\*



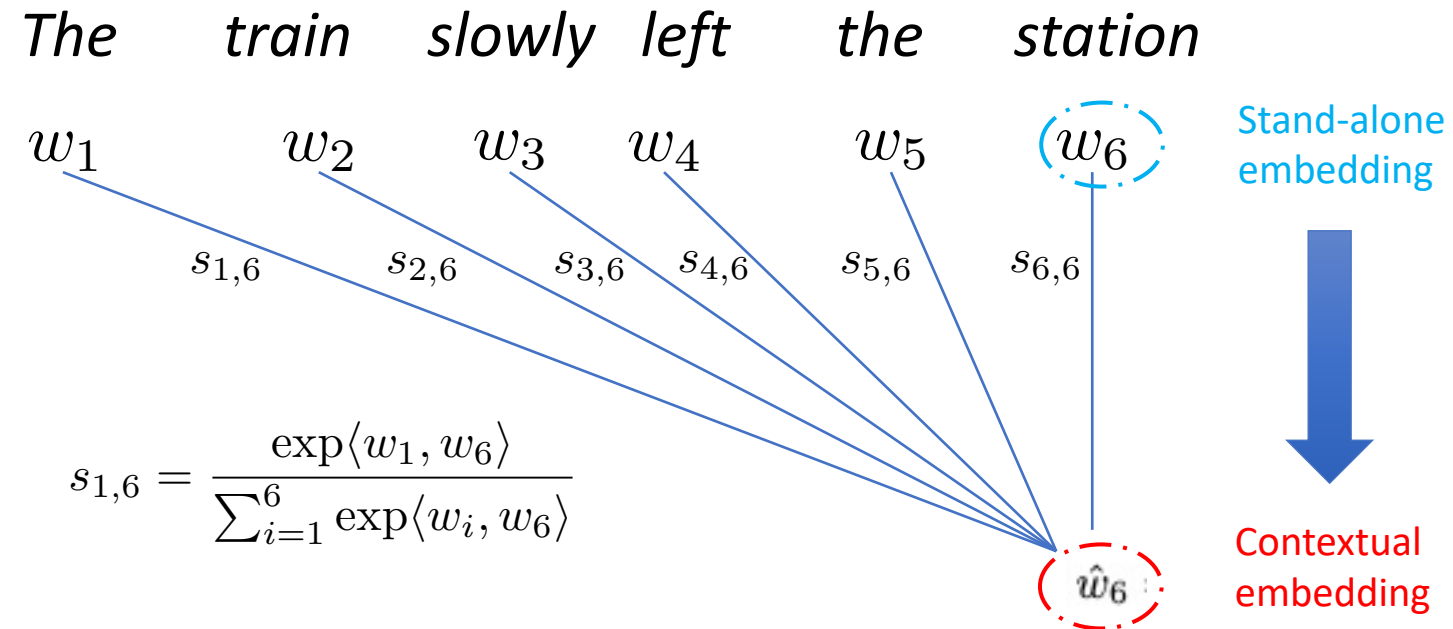
Since dot-products can be negative, we exponentiate them and then normalize (remember softmax?)

\* non-negative, and summing to 1.0

# Normalized attention weights



# From embedding to *contextual* embedding!



$$\hat{w}_6 = s_{1,6}w_1 + s_{2,6}w_2 + s_{3,6}w_3 + s_{4,6}w_4 + s_{5,6}w_5 + s_{6,6}w_6$$

By choosing weights in this manner, the embedding of a word moves closer to the embeddings of the other words in the current context, in proportion to how related they are



- The word 'station' has many contexts.



By choosing weights in this manner, the embedding of a word moves closer to the embeddings of the other words in the current context, in proportion to how related they are



- The word 'station' has many contexts.
  - In the current context, 'train' is closely related to 'station' and therefore exerts a strong “pull” on it

By choosing weights in this manner, the embedding of a word moves closer to the embeddings of the other words in the current context, in proportion to how related they are



- The word 'station' has many contexts.
  - In the current context, 'train' is closely related to 'station' and therefore exerts a strong “pull” on it
  - 'radio' is also related to 'station' but doesn't appear in the current context so (automatically) has zero weight

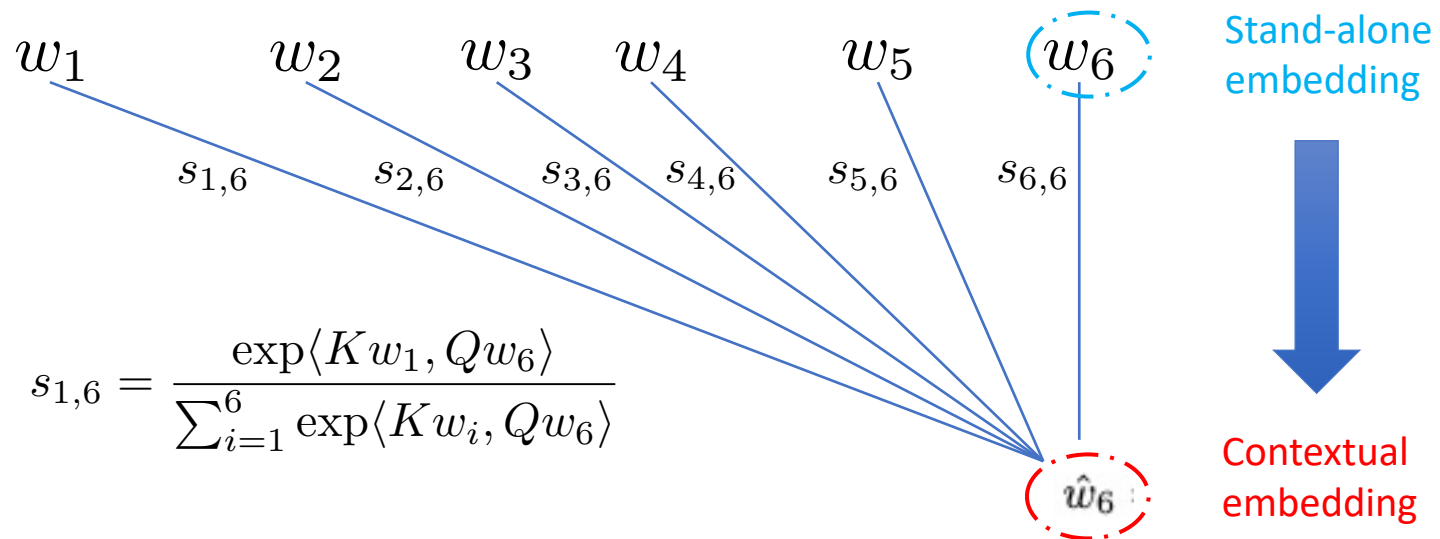
By choosing weights in this manner, the embedding of a word moves closer to the embeddings of the other words in the current context, in proportion to how related they are



- The word 'station' has many contexts.
  - In the current context, 'train' is closely related to 'station' and therefore exerts a strong “pull” on it
  - 'radio' is also related to 'station' but doesn't appear in the current context so (automatically) has zero weight
- By moving station closer to train (equivalently – paying more “attention” to train), we are contextualizing station's embedding to the context of trains, platforms, departures, etc.

# From embedding to *contextual* embedding!

The train slowly left the station



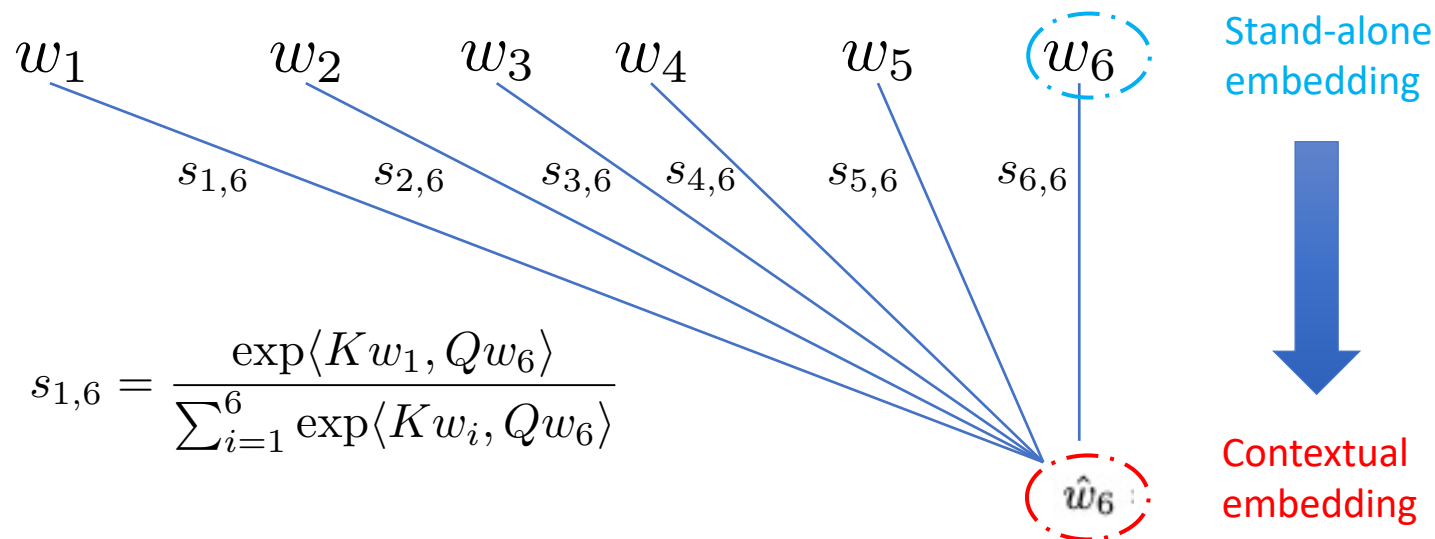
$$s_{1,6} = \frac{\exp\langle K w_1, Q w_6 \rangle}{\sum_{i=1}^6 \exp\langle K w_i, Q w_6 \rangle}$$

$$\hat{w}_6 = s_{1,6}w_1 + s_{2,6}w_2 + s_{3,6}w_3 + s_{4,6}w_4 + s_{5,6}w_5 + s_{6,6}w_6$$

Make it more flexible

# From embedding to *contextual* embedding!

The train slowly left the station



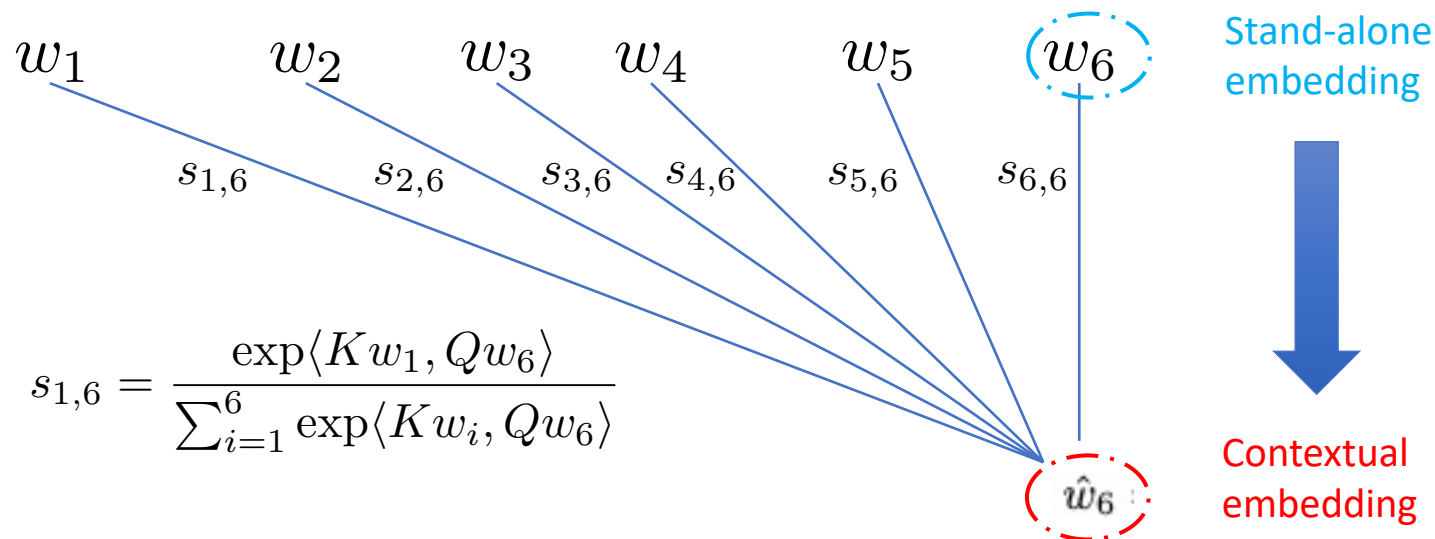
Make it more flexible:  
- e.g., in K=transportation  
subspace, Q=location  
subspace, how closely two  
words are related?

$$s_{1,6} = \frac{\exp\langle K w_1, Q w_6 \rangle}{\sum_{i=1}^6 \exp\langle K w_i, Q w_6 \rangle}$$

$$\hat{w}_6 = s_{1,6}w_1 + s_{2,6}w_2 + s_{3,6}w_3 + s_{4,6}w_4 + s_{5,6}w_5 + s_{6,6}w_6$$

# From embedding to *contextual* embedding!

*The train slowly left the station*



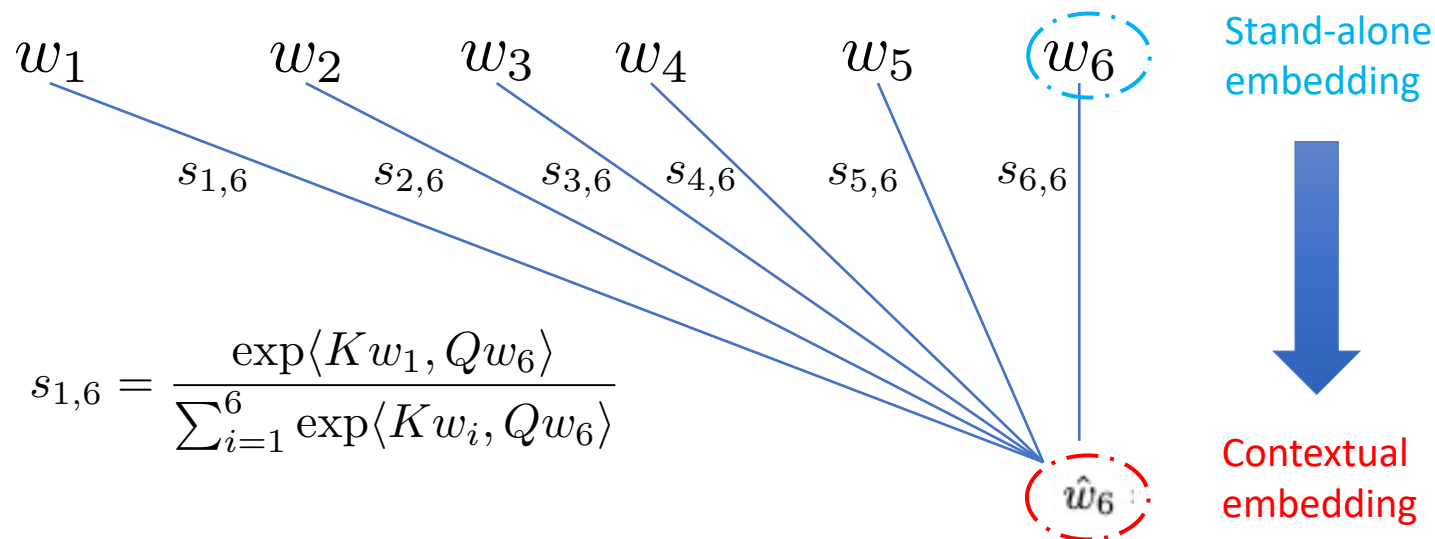
Make it more flexible:  
- e.g., in K=transportation  
subspace, Q=location  
subspace, how closely two  
words are related?

$$s_{1,6} = \frac{\exp\langle K w_1, Q w_6 \rangle}{\sum_{i=1}^6 \exp\langle K w_i, Q w_6 \rangle}$$

$$\hat{w}_6 = s_{1,6} V w_1 + s_{2,6} V w_2 + s_{3,6} V w_3 + s_{4,6} V w_4 + s_{5,6} V w_5 + s_{6,6} V w_6$$

# From embedding to *contextual* embedding!

The train slowly left the station



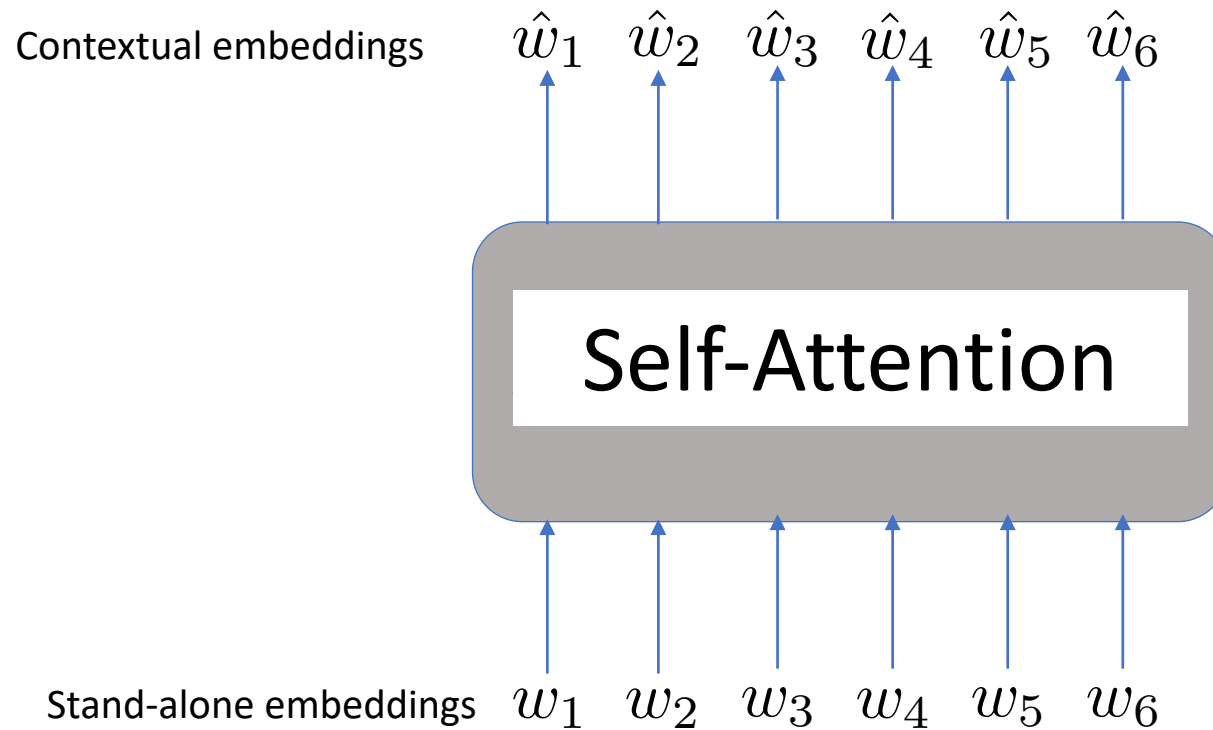
Make it more flexible:  
- e.g., in K=transportation  
subspace, Q=location  
subspace, how closely two  
words are related?

$$s_{1,6} = \frac{\exp\langle K w_1, Q w_6 \rangle}{\sum_{i=1}^6 \exp\langle K w_i, Q w_6 \rangle}$$

$$\hat{w}_6 = s_{1,6} V w_1 + s_{2,6} V w_2 + s_{3,6} V w_3 + s_{4,6} V w_4 + s_{5,6} V w_5 + s_{6,6} V w_6$$

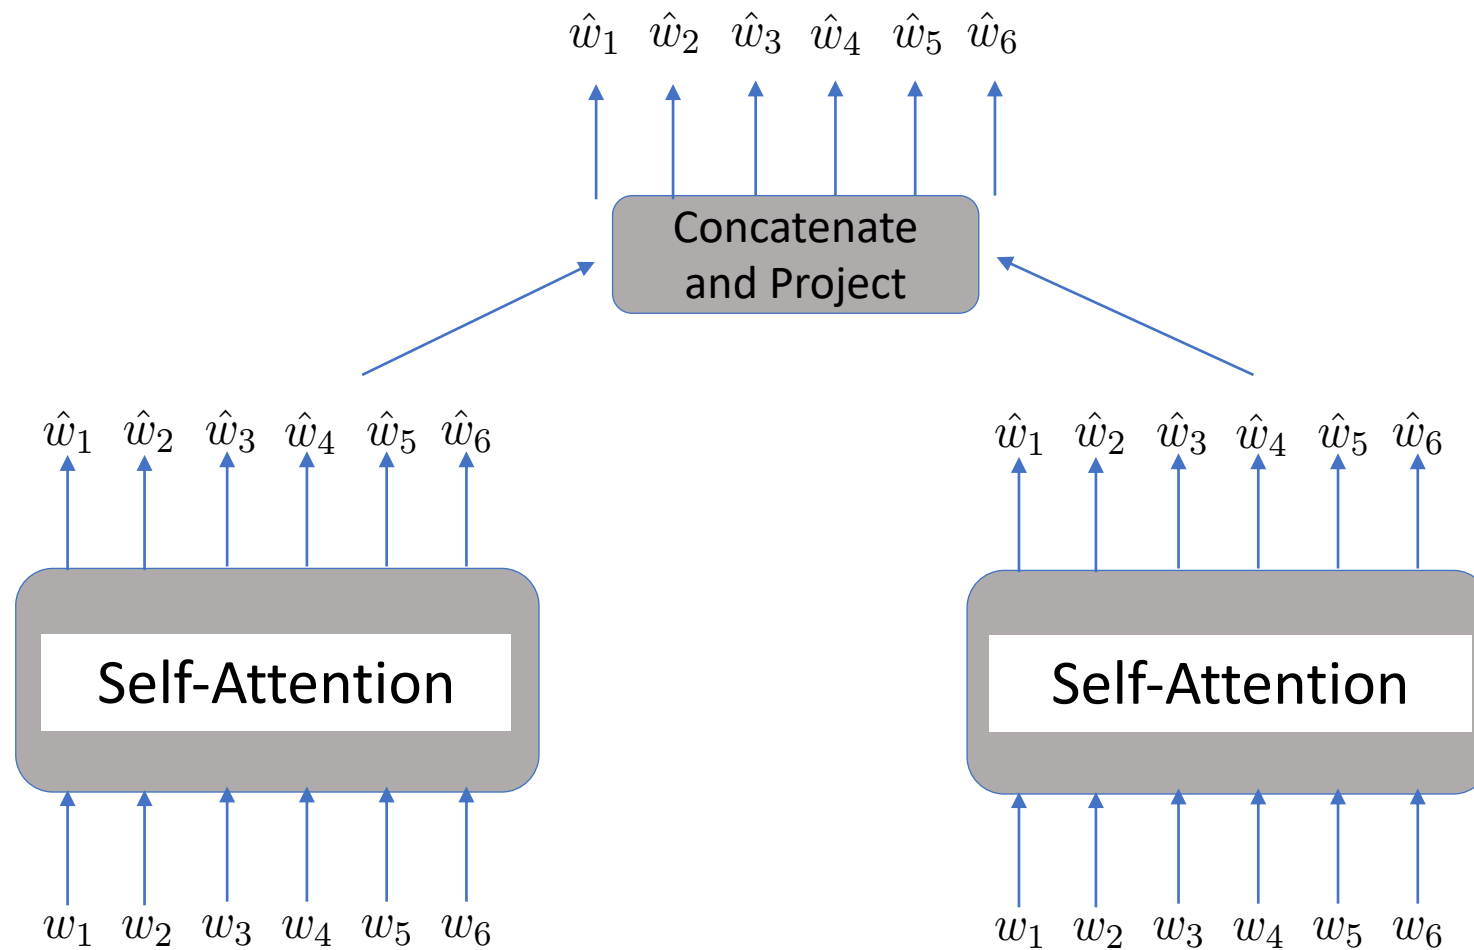
Key (K), Query (Q), Value (V) projection are all tunable parameters:  $K, Q, V \in R^{d \times d_{model}}$

This operation is referred to as a 'Self Attention' layer and can be done very efficiently with matrix operations

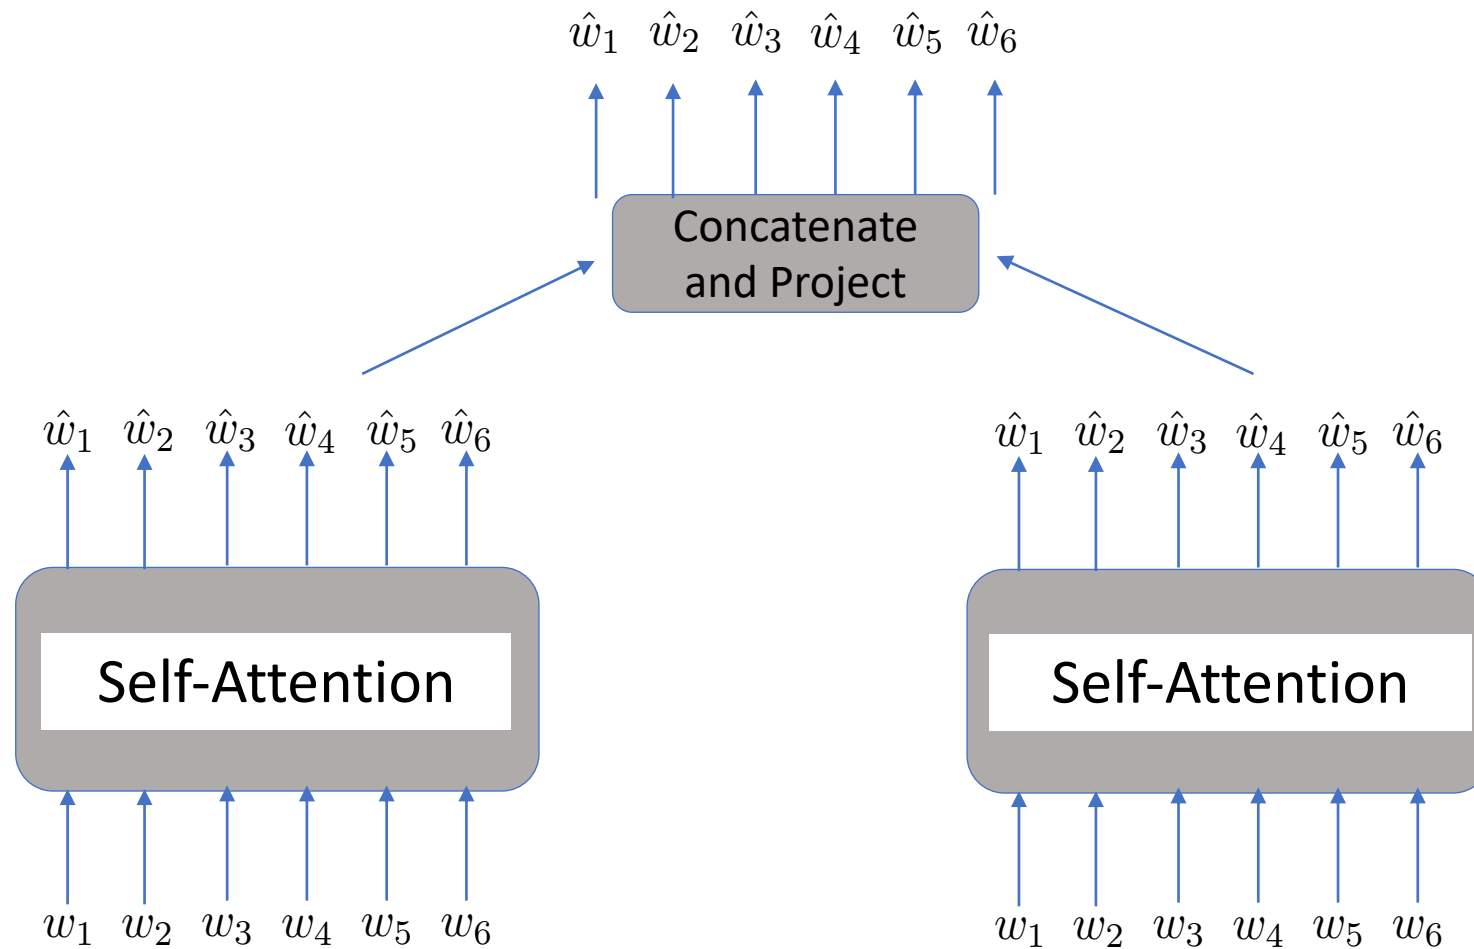




# Key Tweak: Multi\*-Head Attention

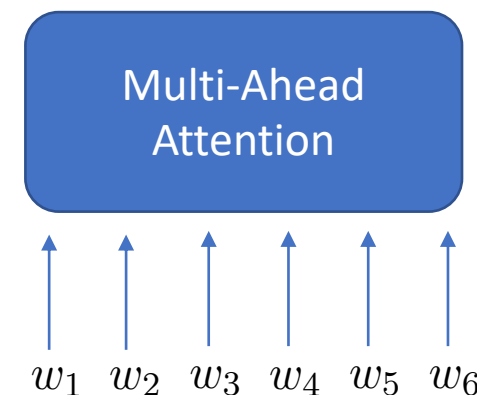
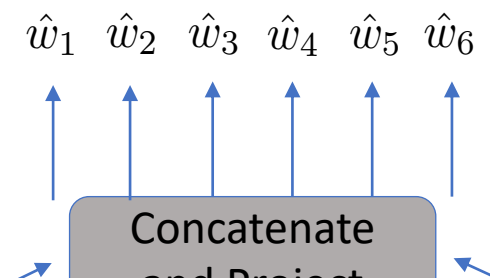


# Key Tweak: Different attention 'heads' learn different patterns

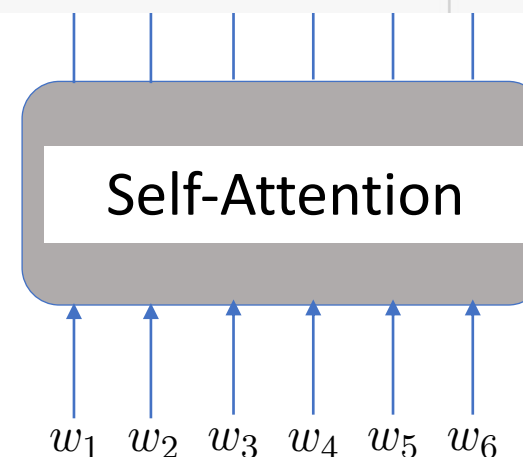
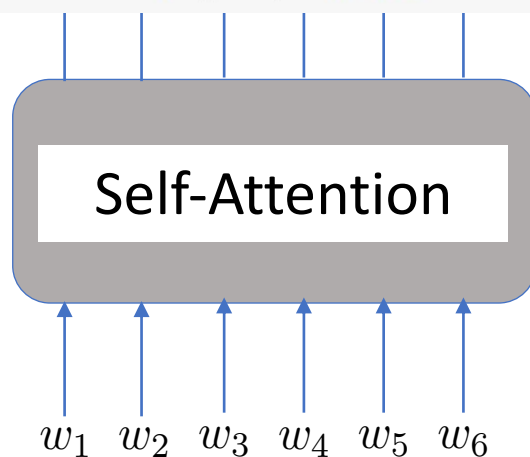


# Key Tweak: Different attention 'heads' learn different patterns

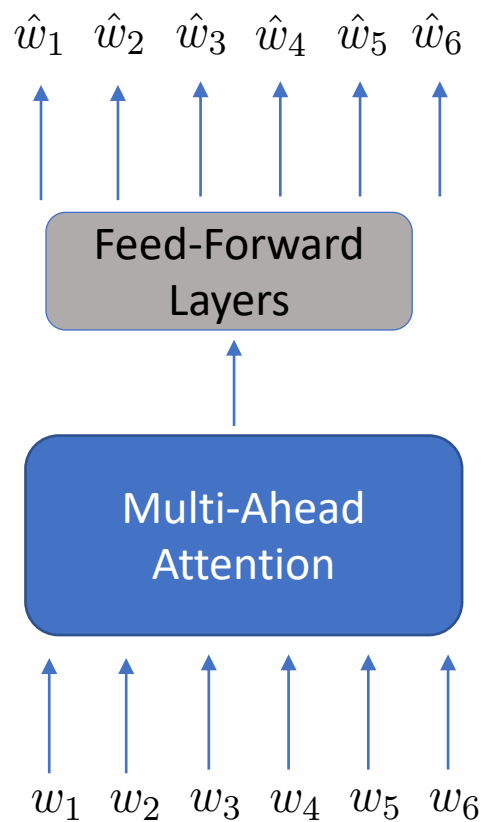
Combine together..



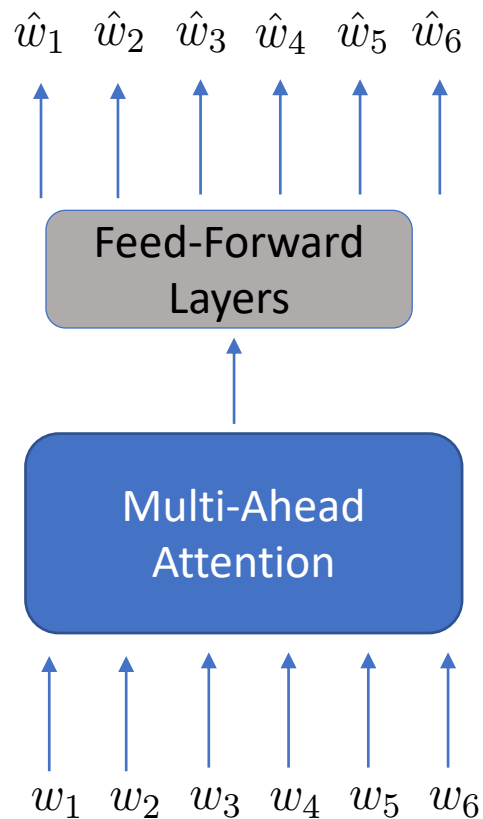
```
def forward(self, x):  
    ### dimension of x is [batch_size, sequence_length, embedding_dim]  
    self.attention = nn.MultiheadAttention(embedding_dim, num_heads, dropout=dropout, batch_first=True)  
    return self.attention(x, x, x)[0]
```



# Key Tweak: Inject some non-linearity with feed-forward layers at the end

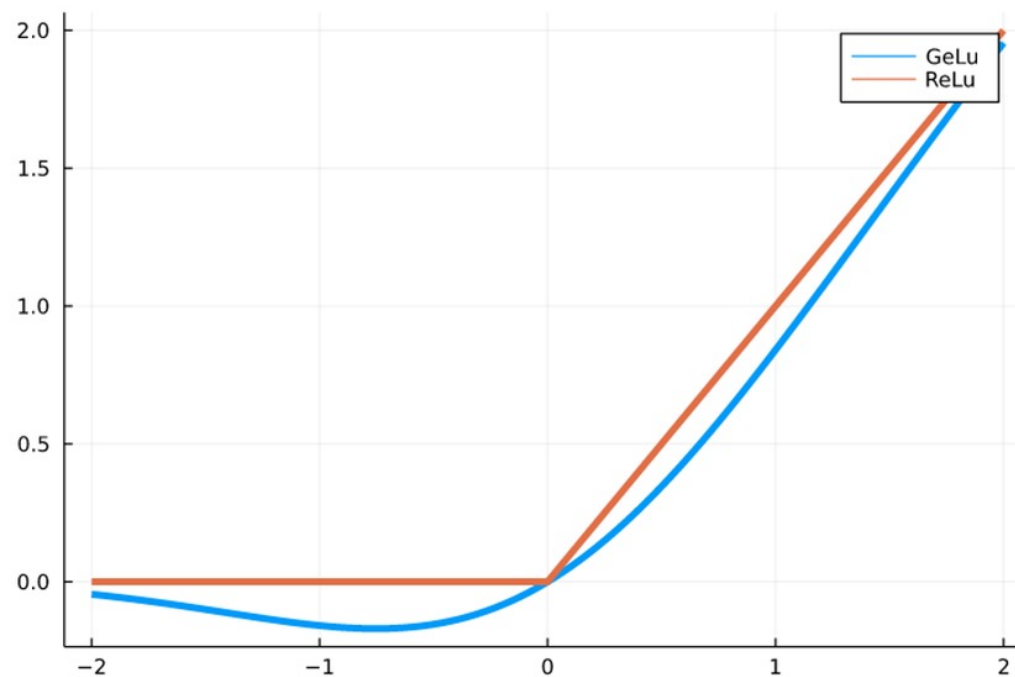


# Key Tweak: Inject some non-linearity with feed-forward layers at the end



```
self.feed_forward = nn.Sequential(  
    nn.Linear(embedding_dim, 4 * embedding_dim),  
    nn.GELU(),  
    nn.Linear(4 * embedding_dim, embedding_dim),  
)
```

$$GELU(x) = xP(X \leq x), X \sim N(0, 1)$$

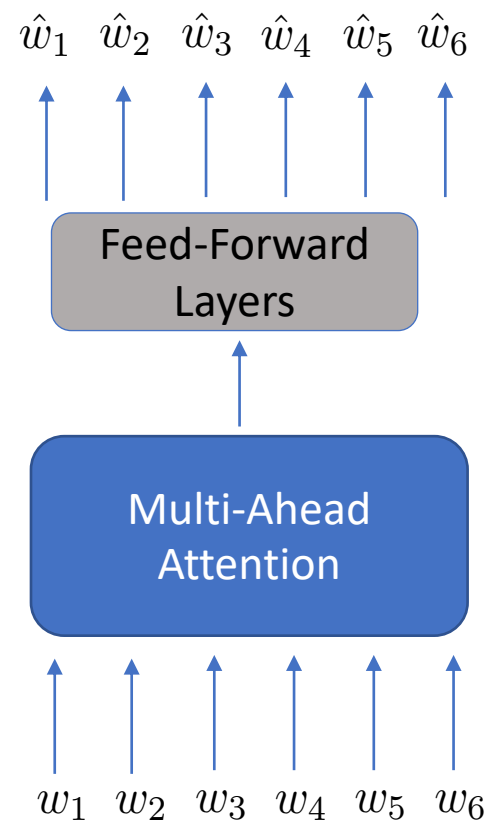


# The story so far

*End with contextual embeddings*



*Start with random embeddings*

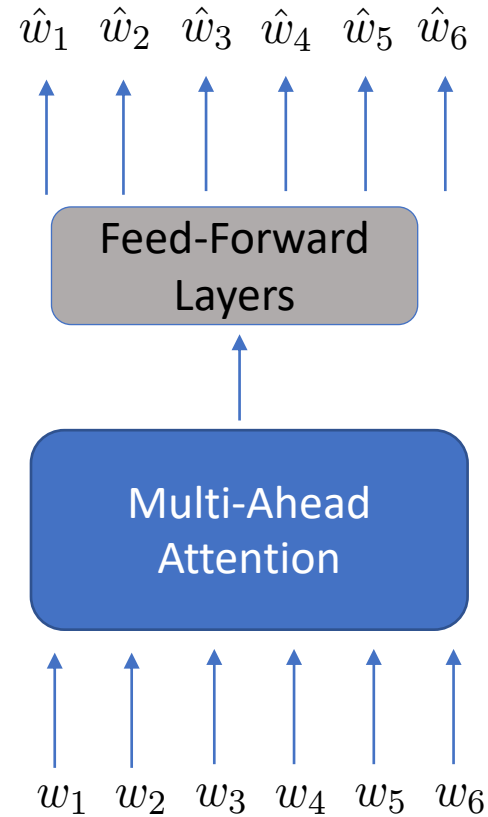


# The story so far

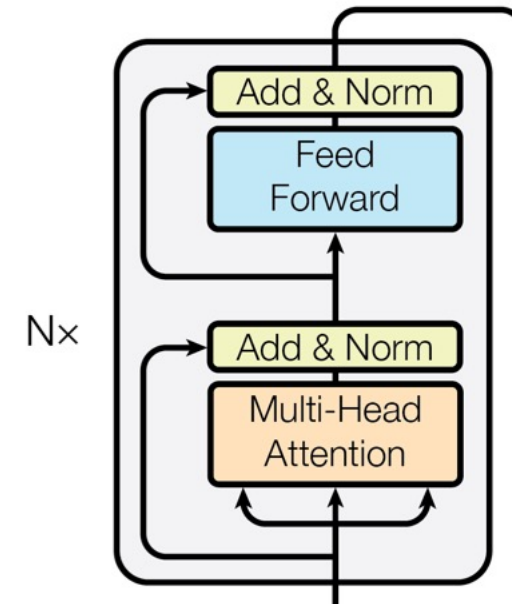
*End with contextual embeddings*



*Start with random embeddings*



Transformer Layer



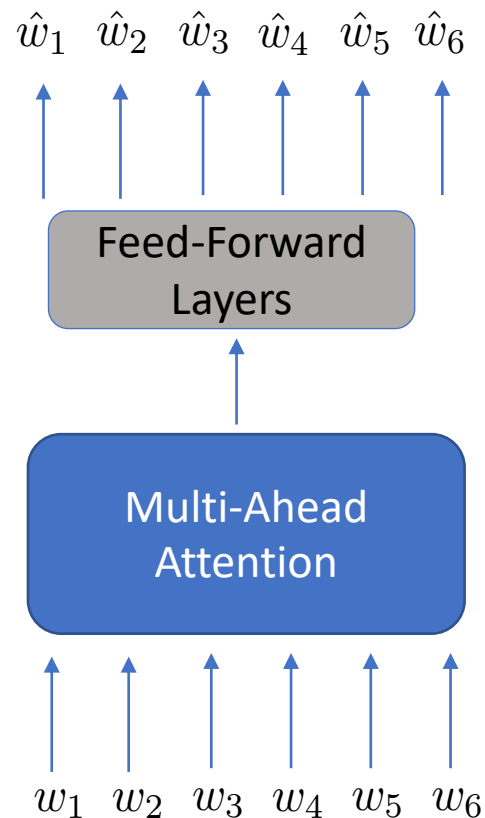
Attention is all you need

# The story so far

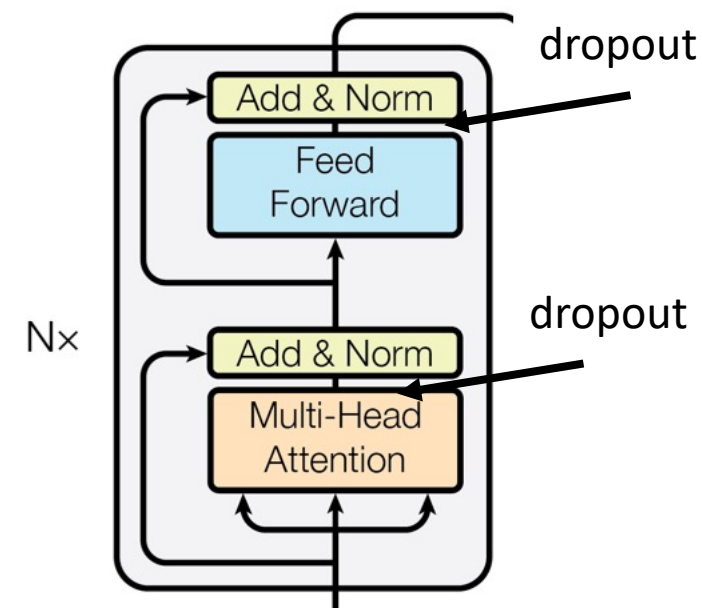
*End with contextual embeddings*



*Start with random embeddings*



Transformer Layer



Attention is all you need

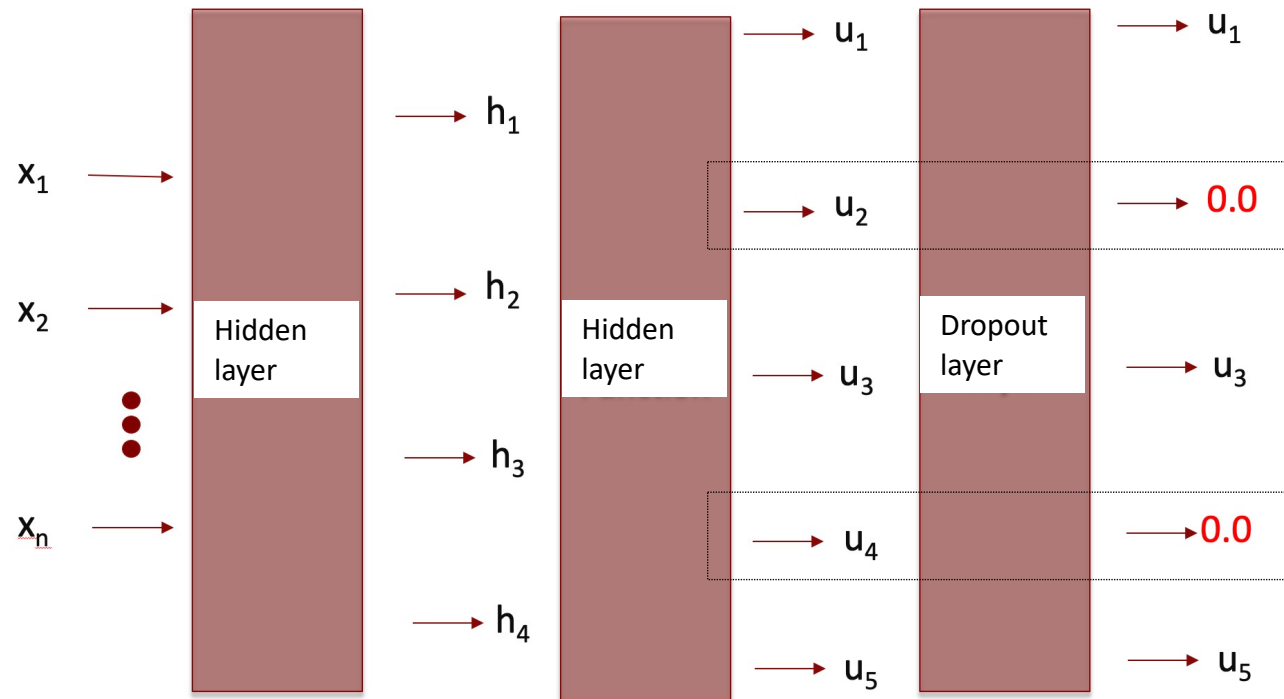


# Three regularization techniques for Deep Learning

- Dropout
- Residual connections (Add)
- Layer normalization (Norm)

# Dropout

In Training, randomly zero out the output from some of the nodes (typically 10% of the nodes) in a hidden layer (implemented as a “dropout layer” in PyTorch)

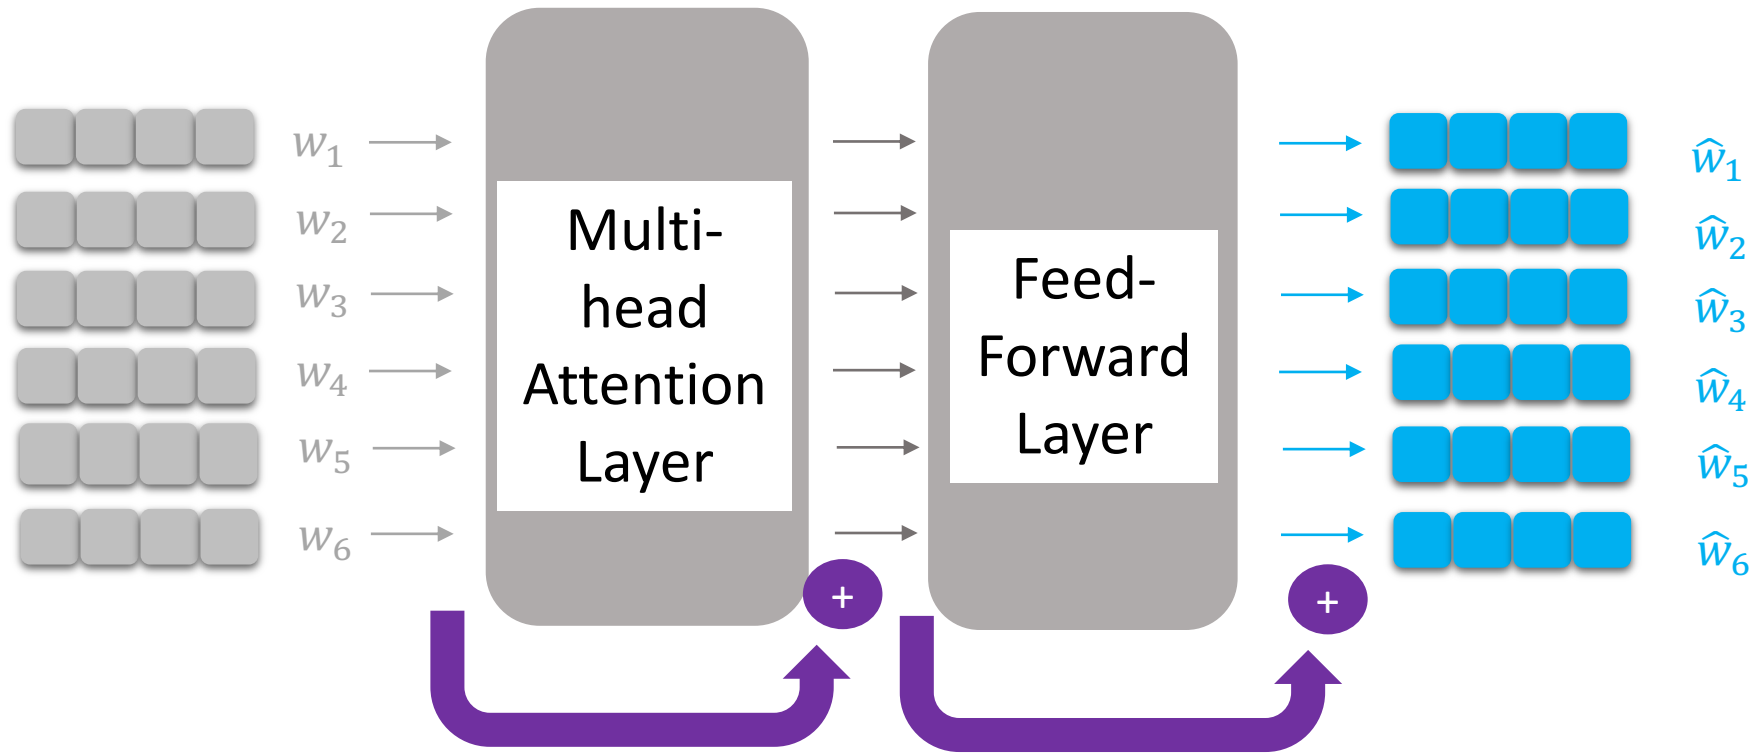


```
self.dropout = nn.Dropout(0.1)  
x = self.dropout(x)
```

# Residual Connection

$$y = \text{Layer}(x) + x$$

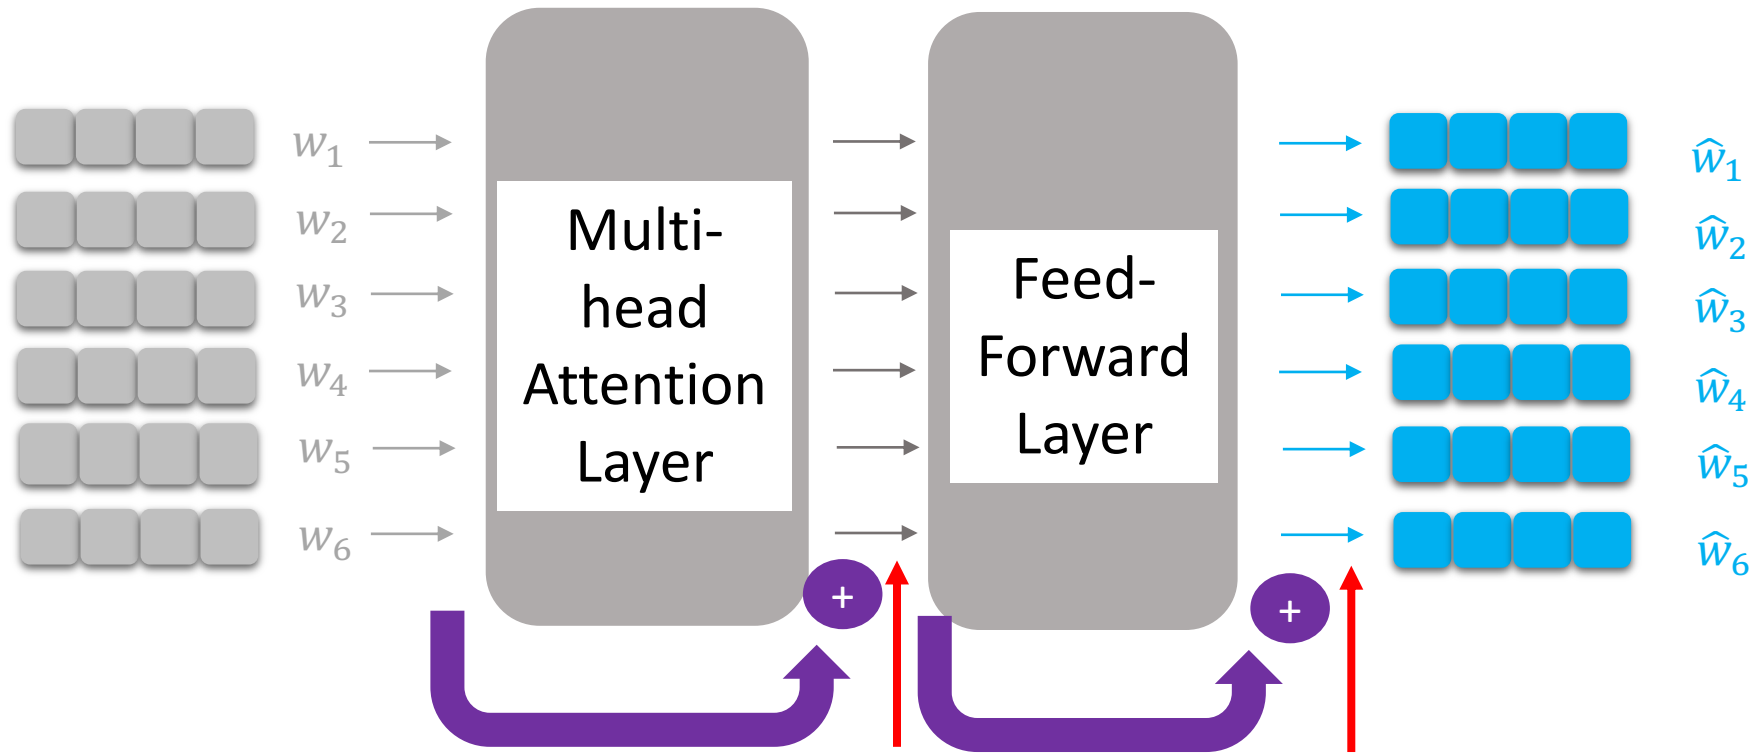
We **sum the input embedding** to the output embedding of the Attention / Feed-Forward Layers. This helps gradients flow better during backpropagation.



# Layer Normalization

$$y = \frac{x - \mathbb{E}[x]}{\sqrt{\text{Var}[x] + \epsilon}} * \gamma + \beta$$

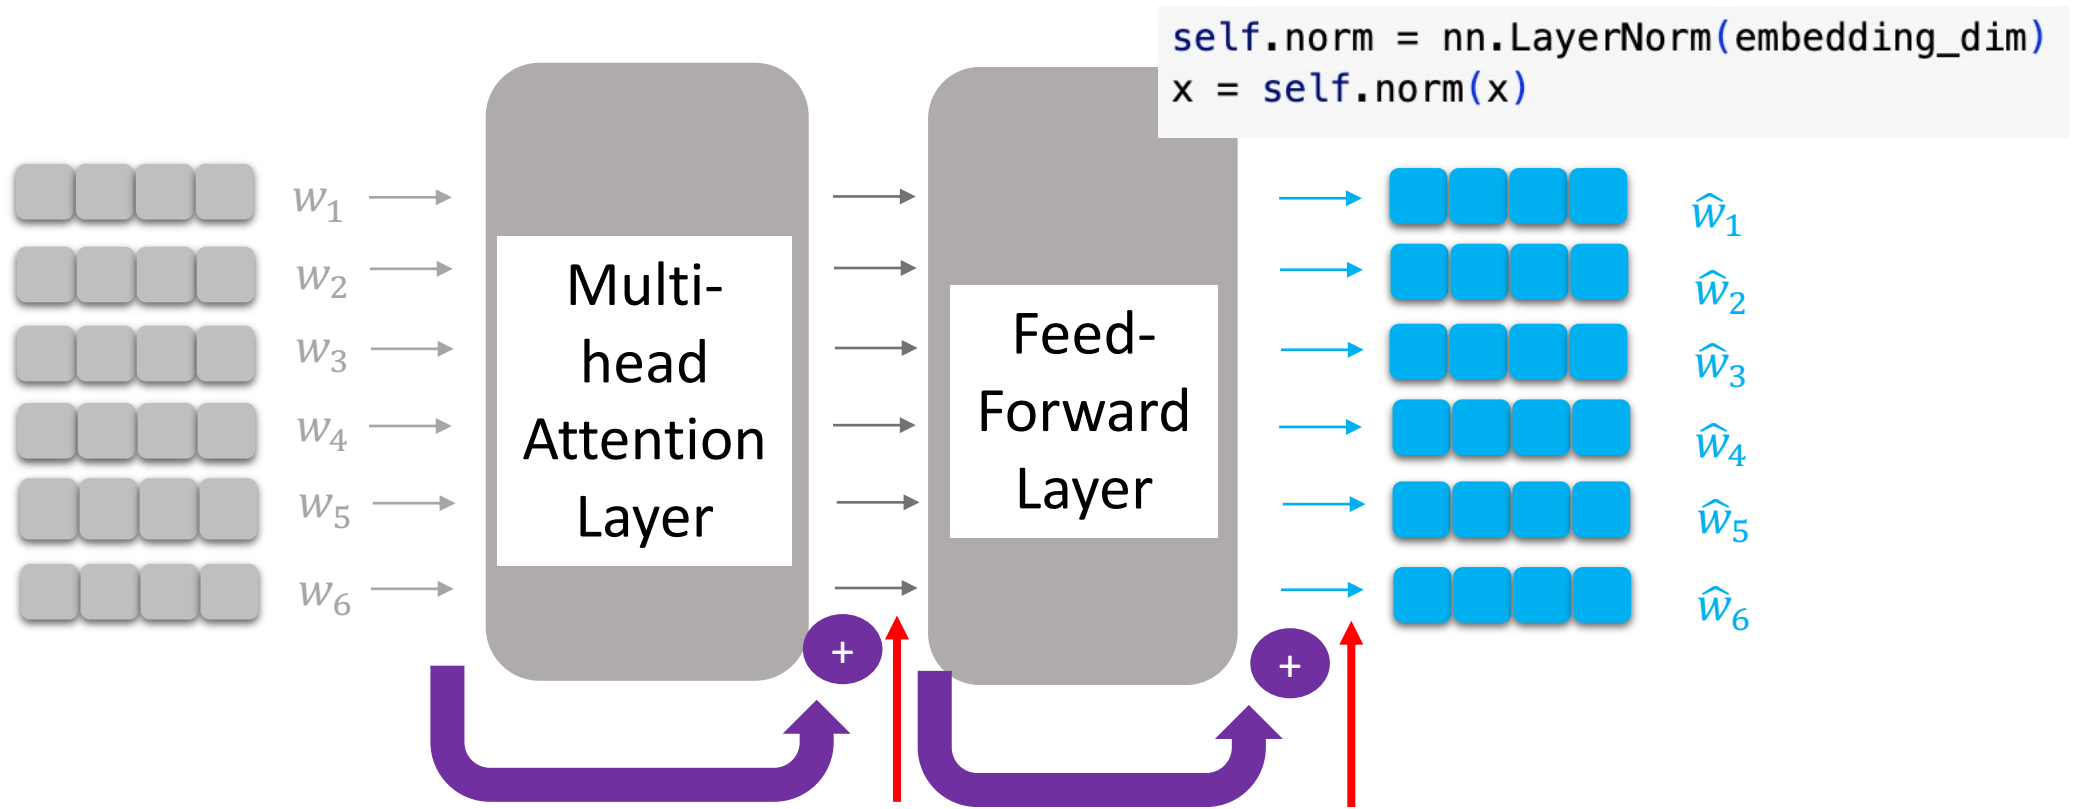
After the Attention / Feed-Forward Layers, we **standardize** (i.e., subtract mean and divide by std) each embedding. This ensures that the weights stay small.



# Layer Normalization

$$y = \frac{x - \mathbb{E}[x]}{\sqrt{\text{Var}[x] + \epsilon}} * \gamma + \beta$$

After the Attention / Feed-Forward Layers, we **standardize** (i.e., subtract mean and divide by std) each embedding. This ensures that the weights stay small.

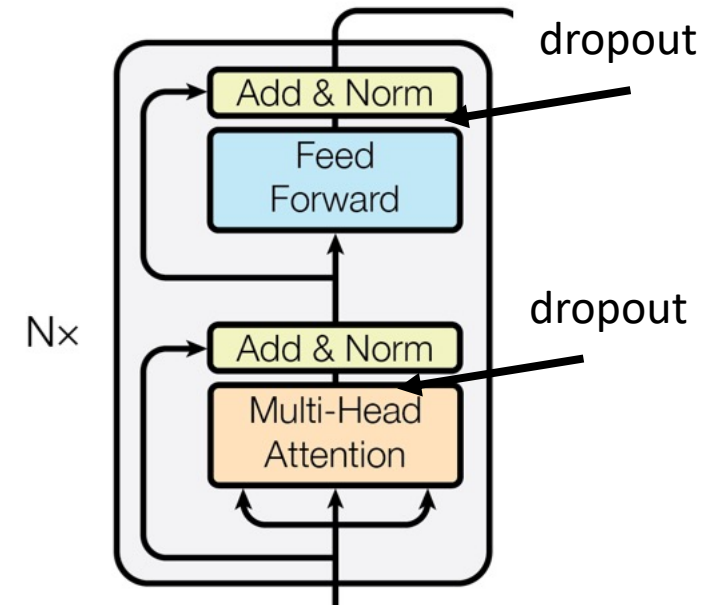


# Taken Together: A Transformer Layer

```
y = MultiheadAttention(x) ## Attention Layer
y = Dropout(y) ## Dropout layer
x = x + y ## Add Residue
x = LayerNorm(x) ## Layer Norm

y = FeedForward(x) ## Feed Forward Layer
y = Dropout(y) ## Dropout layer
x = x + y ## Add Residue
x = LayerNorm(x) ## Layer Norm
```

Pseudo code



Attention is all you need

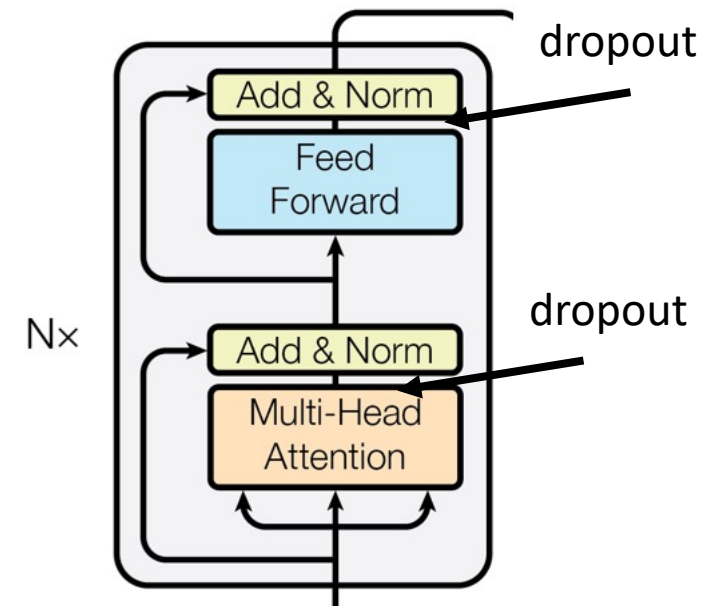
# Taken Together: A Transformer Layer

What is dimension of x?

```
y = MultiheadAttention(x) ## Attention Layer
y = Dropout(y) ## Dropout layer
x = x + y ## Add Residue
x = LayerNorm(x) ## Layer Norm

y = FeedForward(x) ## Feed Forward Layer
y = Dropout(y) ## Dropout layer
x = x + y ## Add Residue
x = LayerNorm(x) ## Layer Norm
```

Pseudo code



Attention is all you need

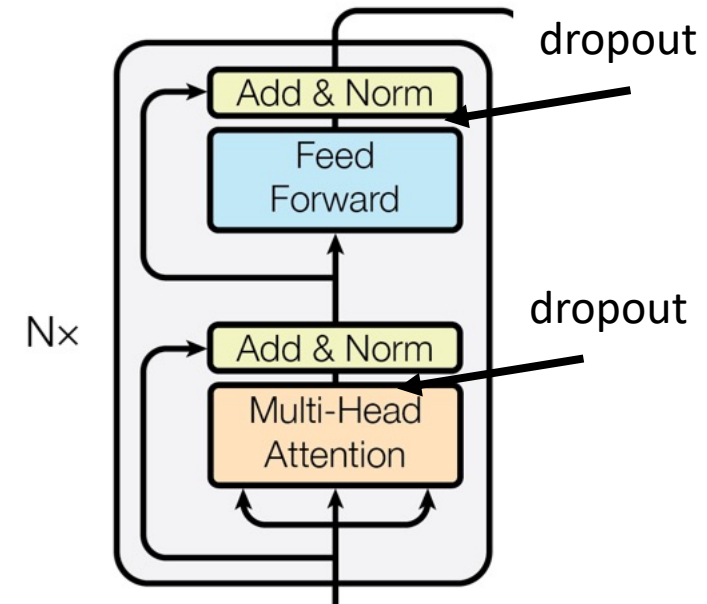
# Taken Together: A Transformer Layer

What is dimension of x? [batch\_size, seq\_len, embedding\_dim]

```
y = MultiheadAttention(x) ## Attention Layer
y = Dropout(y) ## Dropout layer
x = x + y ## Add Residue
x = LayerNorm(x) ## Layer Norm

y = FeedForward(x) ## Feed Forward Layer
y = Dropout(y) ## Dropout layer
x = x + y ## Add Residue
x = LayerNorm(x) ## Layer Norm
```

Pseudo code



Attention is all you need



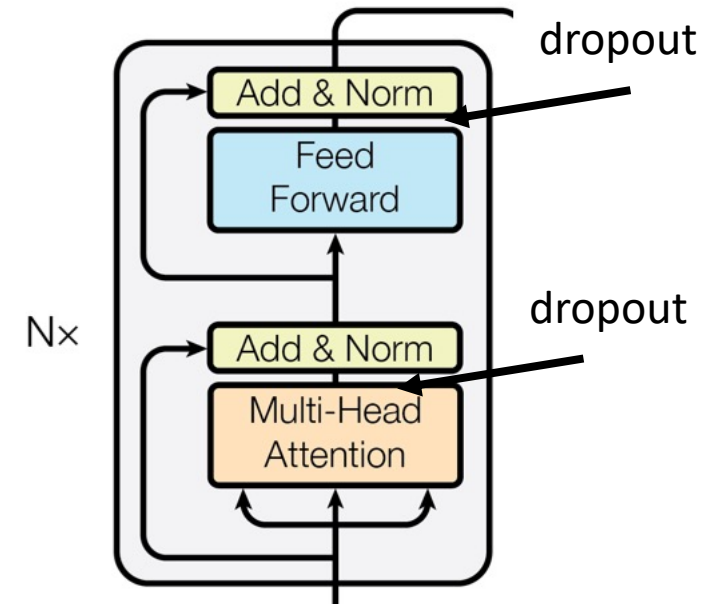
# Taken Together: A Transformer Layer

What is dimension of x? [batch\_size, seq\_len, embedding\_dim]

```
y = MultiheadAttention(x) ## Attention Layer
y = Dropout(y) ## Dropout layer
x = x + y ## Add Residue
x = LayerNorm(x) ## Layer Norm

y = FeedForward(x) ## Feed Forward Layer
y = Dropout(y) ## Dropout layer
x = x + y ## Add Residue
x = LayerNorm(x) ## Layer Norm
```

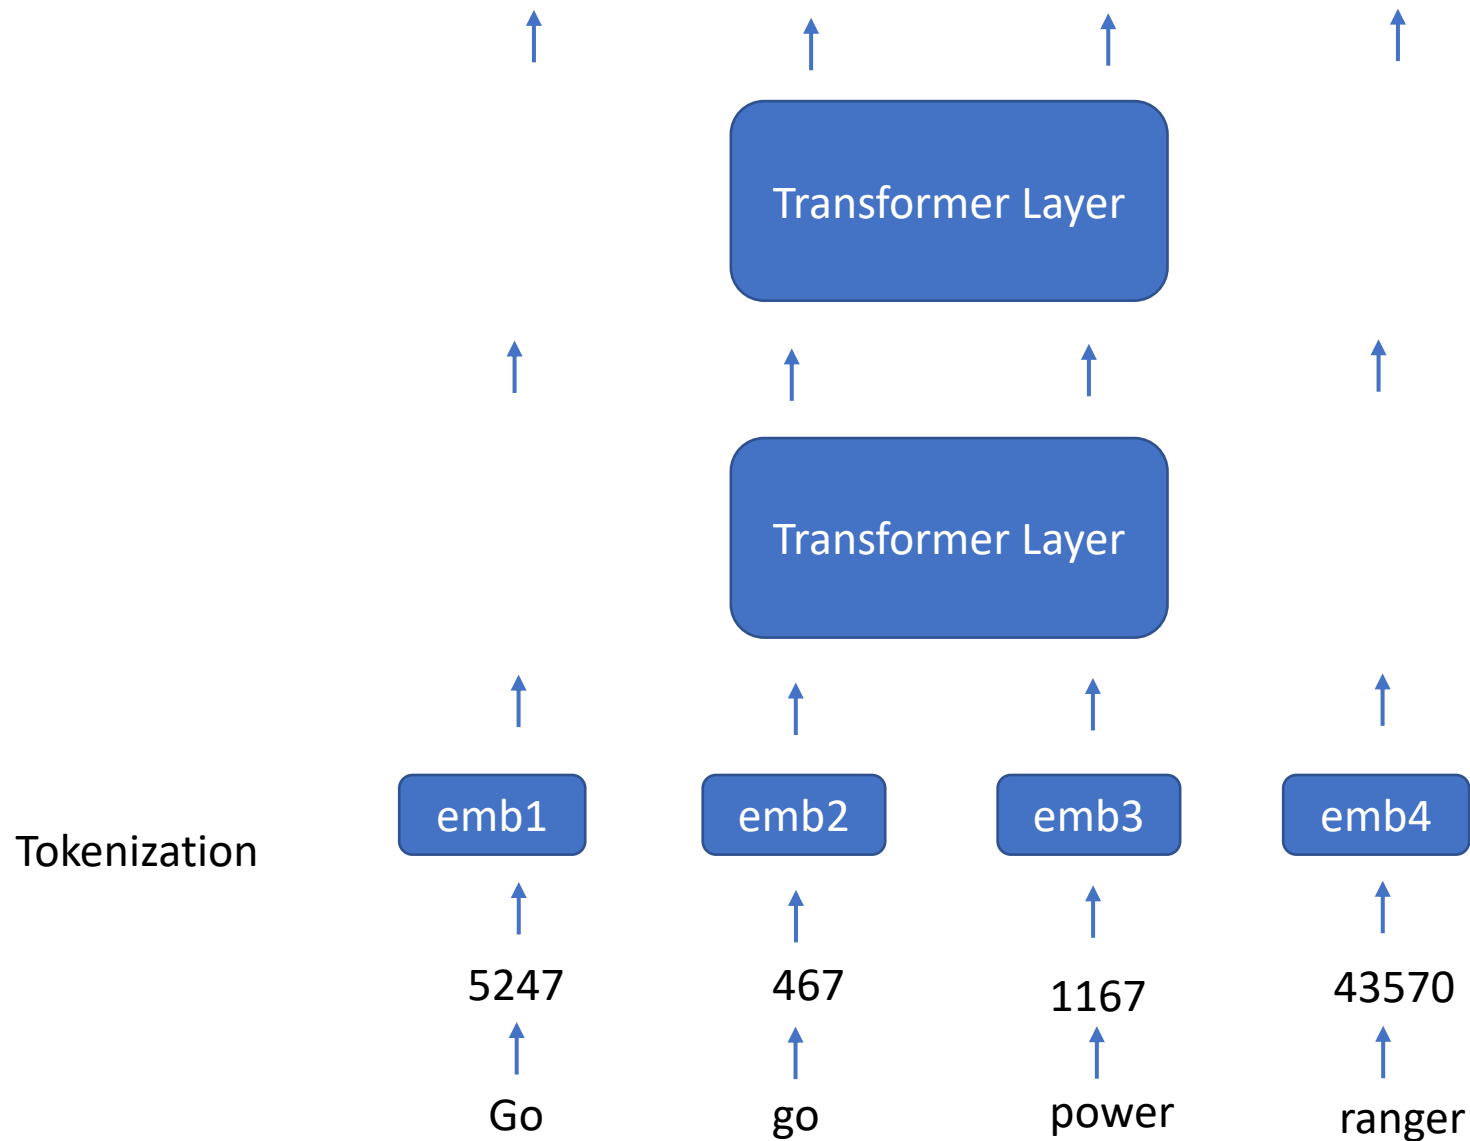
Pseudo code



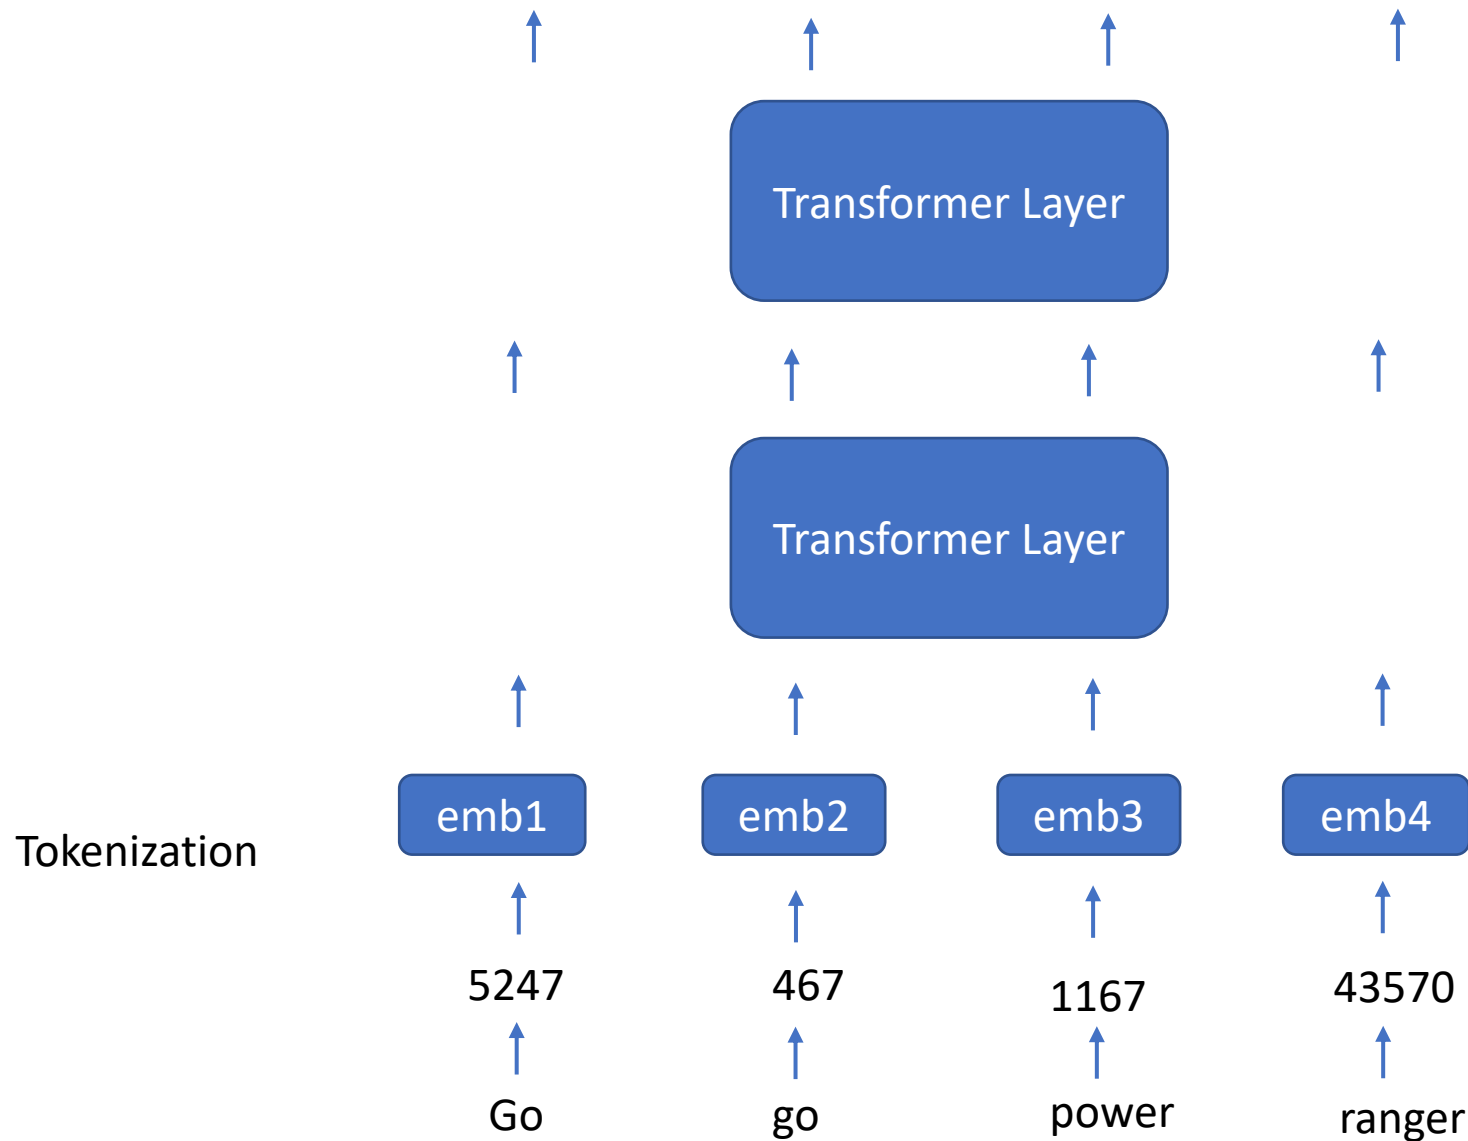
Attention is all you need

We can have multiple Transformer Layers stacked

# Does everything look good?

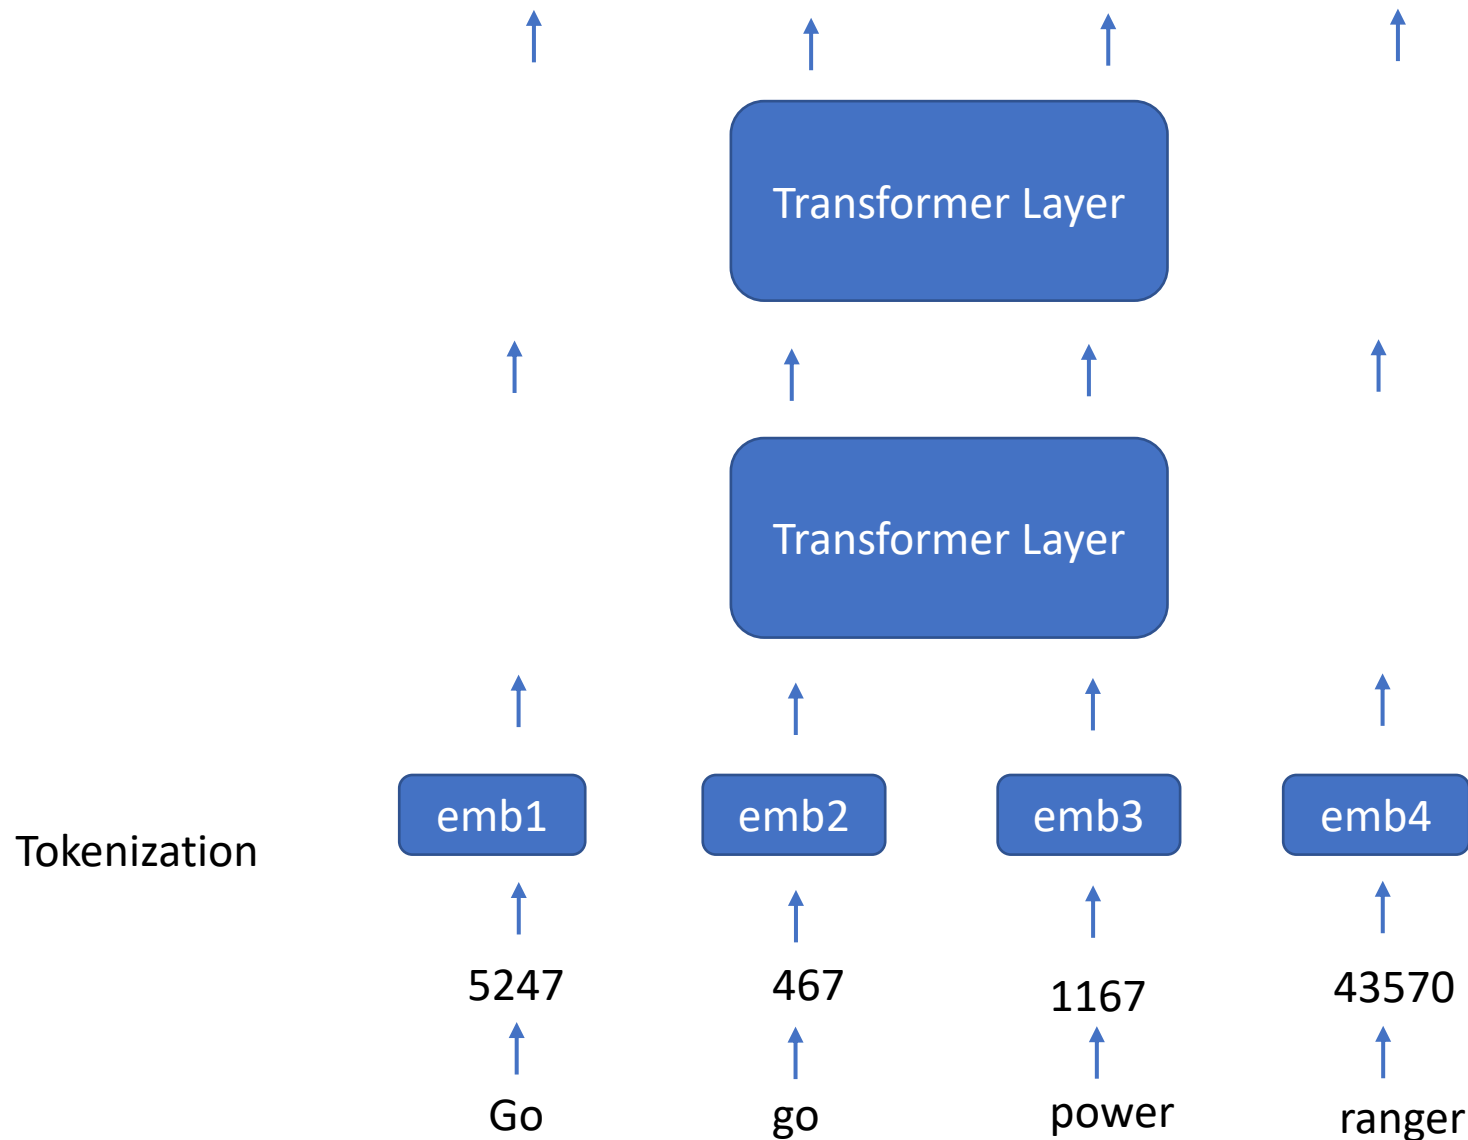


# Does everything look good?



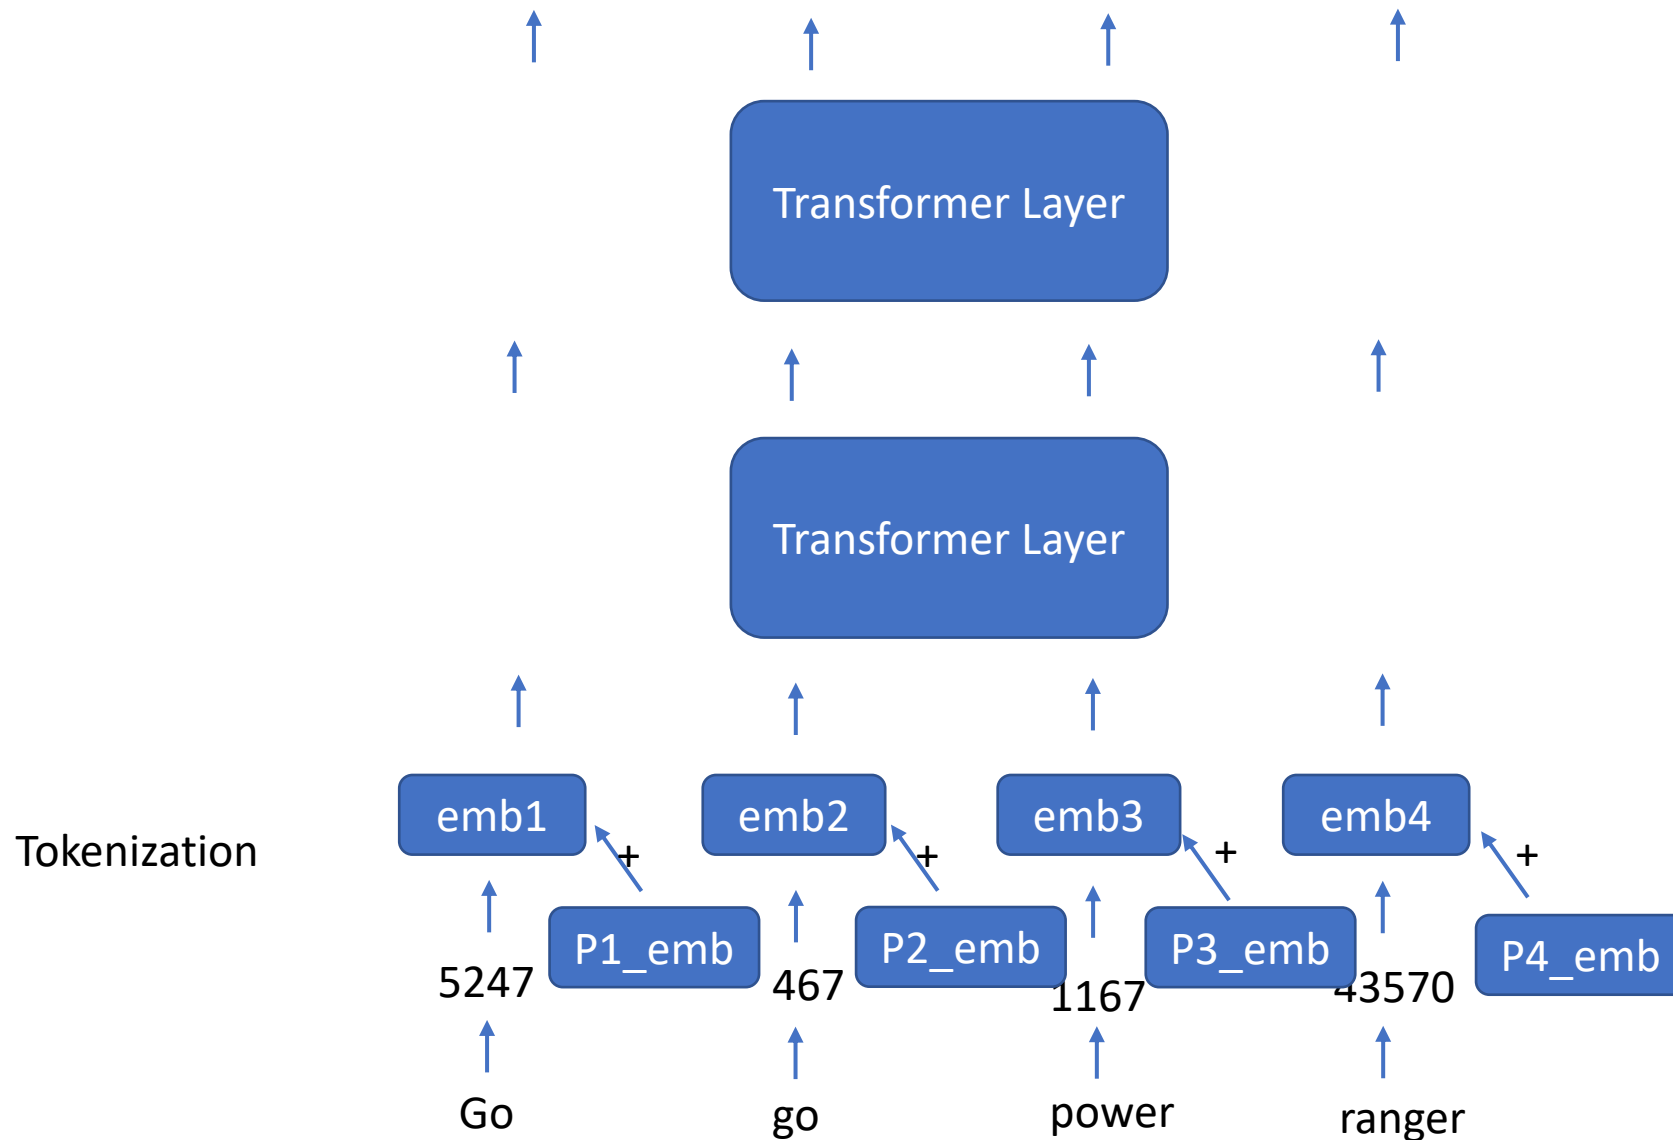
Order information of words  
are missing

# Does everything look good?



Let's directly take positions as inputs and learn a position embedding!

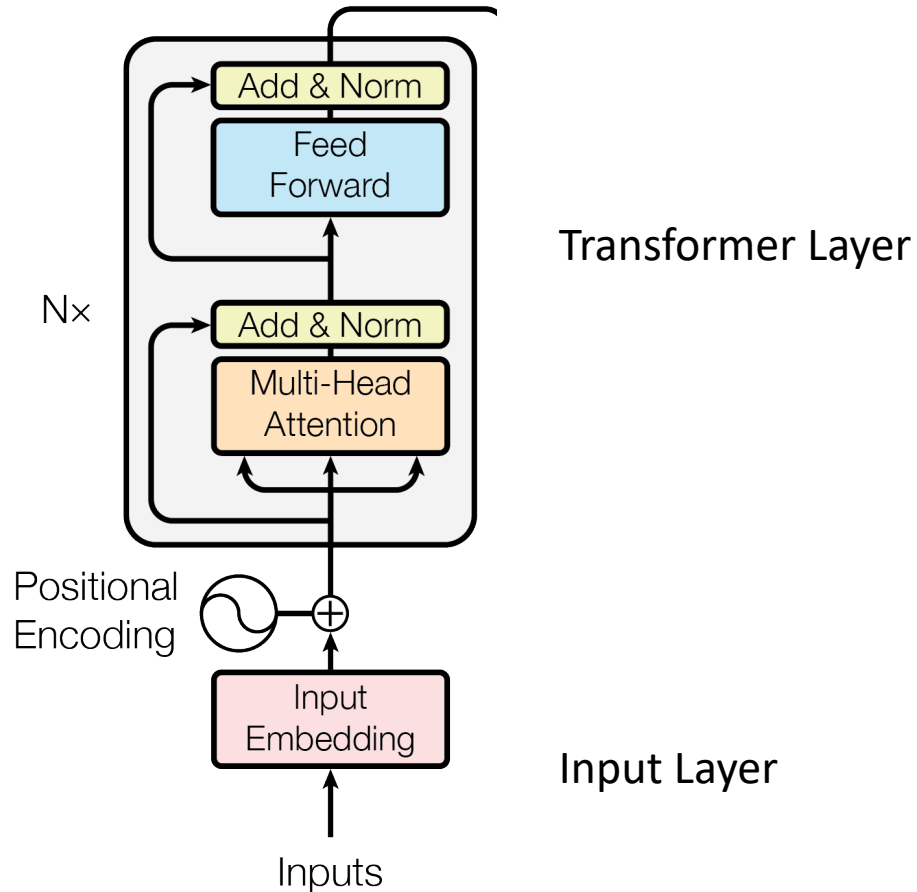
# Does everything look good?



Let's directly take positions as inputs and learn a position embedding!

# This is called a Transformer (Encoder)!

What is dimension of  $x$ ? [batch\_size, seq\_len, embedding\_dim]



```
y = MultiheadAttention(x) ## Attention Layer  
y = Dropout(y) ## Dropout layer  
x = x + y ## Add Residue  
x = LayerNorm(x) ## Layer Norm
```

```
y = FeedForward(x) ## Feed Forward Layer  
y = Dropout(y) ## Dropout layer  
x = x + y ## Add Residue  
x = LayerNorm(x) ## Layer Norm
```

```
x = Embedding(inputs)  
x = x + Embedding(torch.arange(seq_len))  
x = LayerNorm(x)  
x = Dropout(x)
```

<https://arxiv.org/abs/1706.03762>

# Coding Exercise

- Text Classification:

[https://colab.research.google.com/drive/1hWA0Tf4DpeSbhzpb5DP6M\\_KxfvZOxEyw?usp=sharing](https://colab.research.google.com/drive/1hWA0Tf4DpeSbhzpb5DP6M_KxfvZOxEyw?usp=sharing)

```
x = Transformer(inputs)
x = torch.mean(x, dim=1)
x = nn.Linear(embedding_dim, 2)(x)
```

Classification Head

# A better training wrapper

- DataLoader
  - The DataLoader in PyTorch efficiently batches and shuffles data, while also enabling parallel processing and flexible collation, making data loading faster and more scalable for training large models.
- Pytorch\_lightning
  - streamlines the training process by automating loops, logging, and device management, enabling cleaner, more efficient model development.



# Text Generation

A journey of a thousand miles begins with a single step. ---- Laozi (老子)

How would we train a model to generate the next token/word?

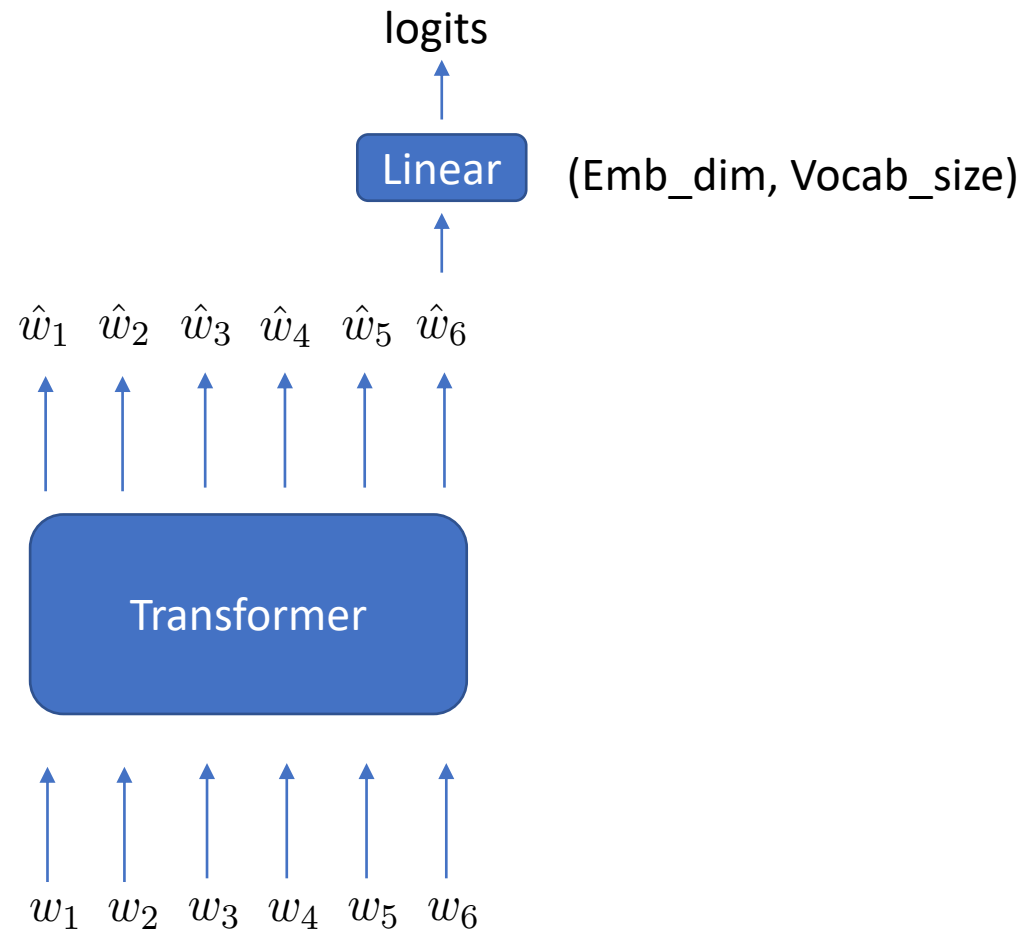
# How would we train a model to generate the next token/word?

- Multi-class classification problem!

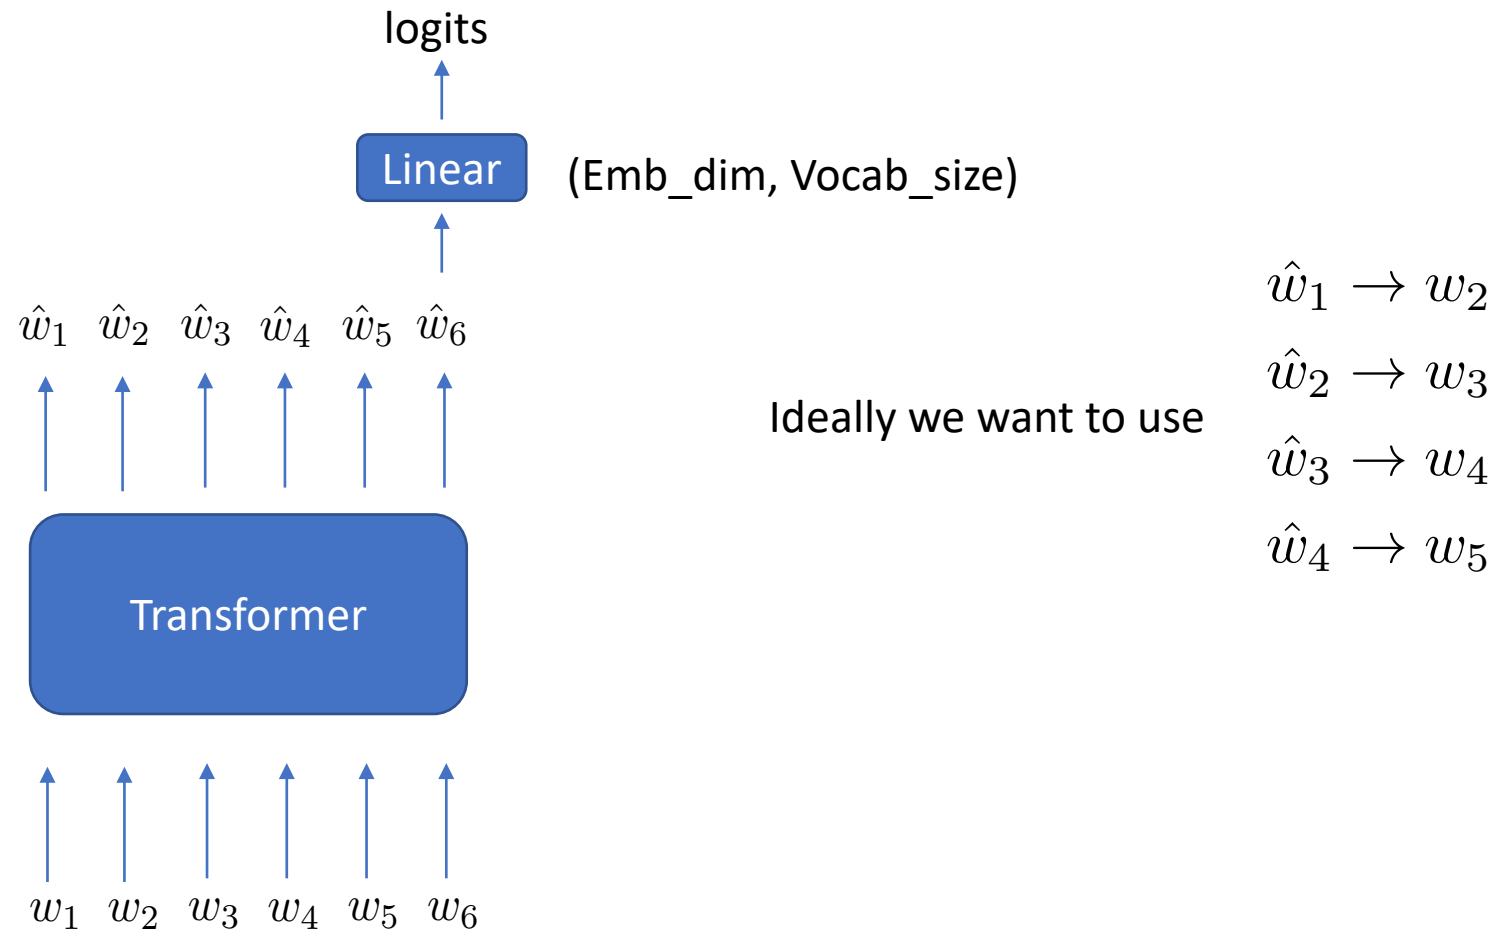
Next Word*	Probability
aardvark	0.0003
...	
rainy	0.3
...	
stormy	0.6
...	
zebra	0.00009

Cross Entropy Loss

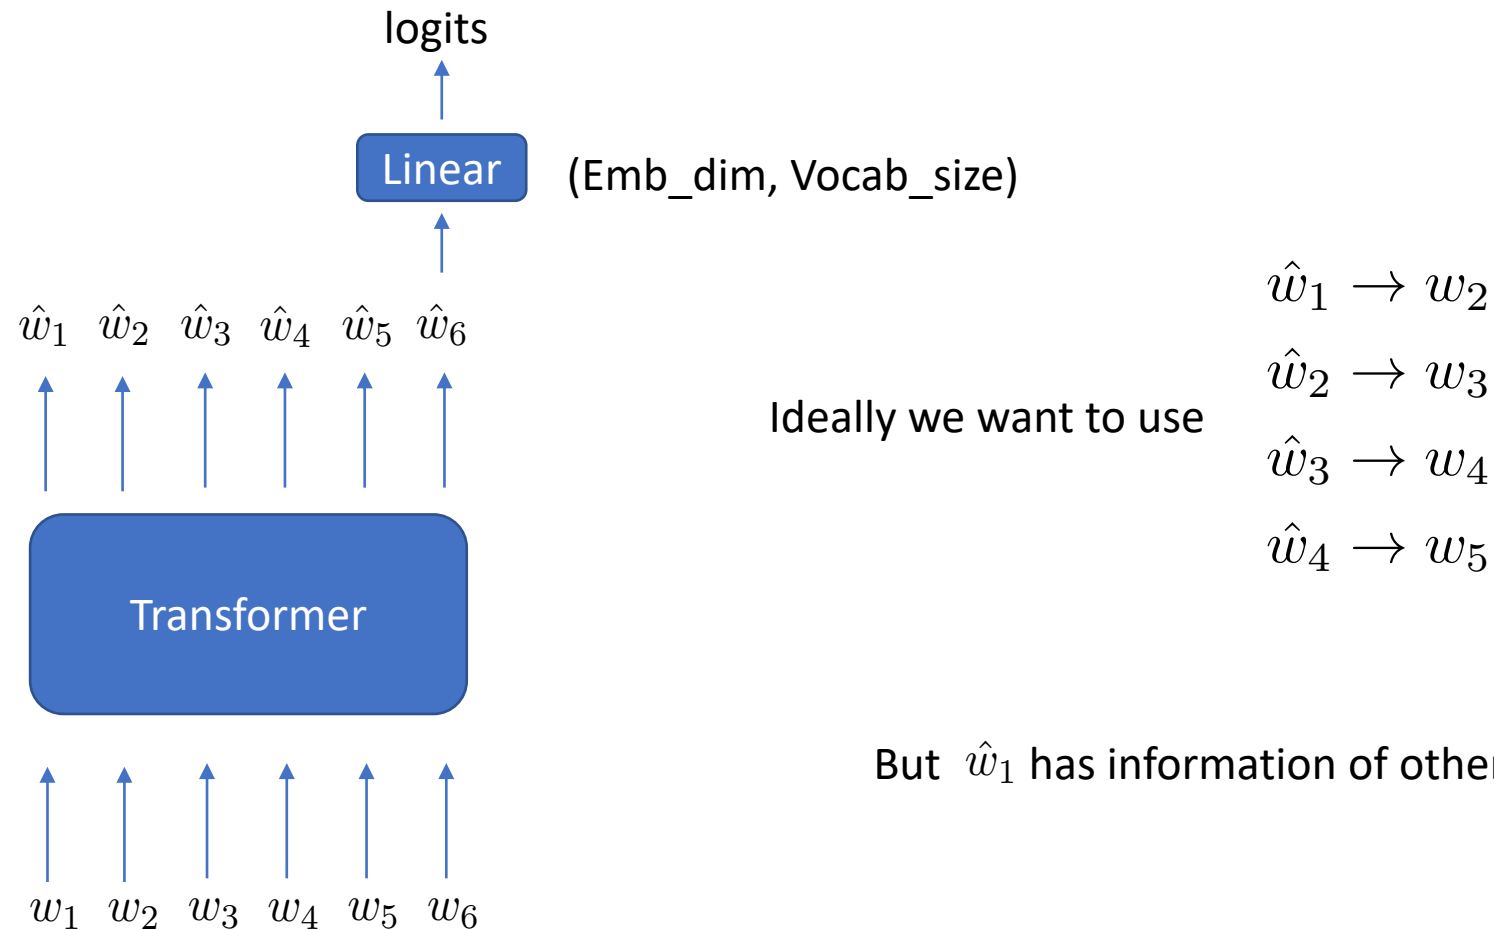
# How would we train a model to generate the next token/word?



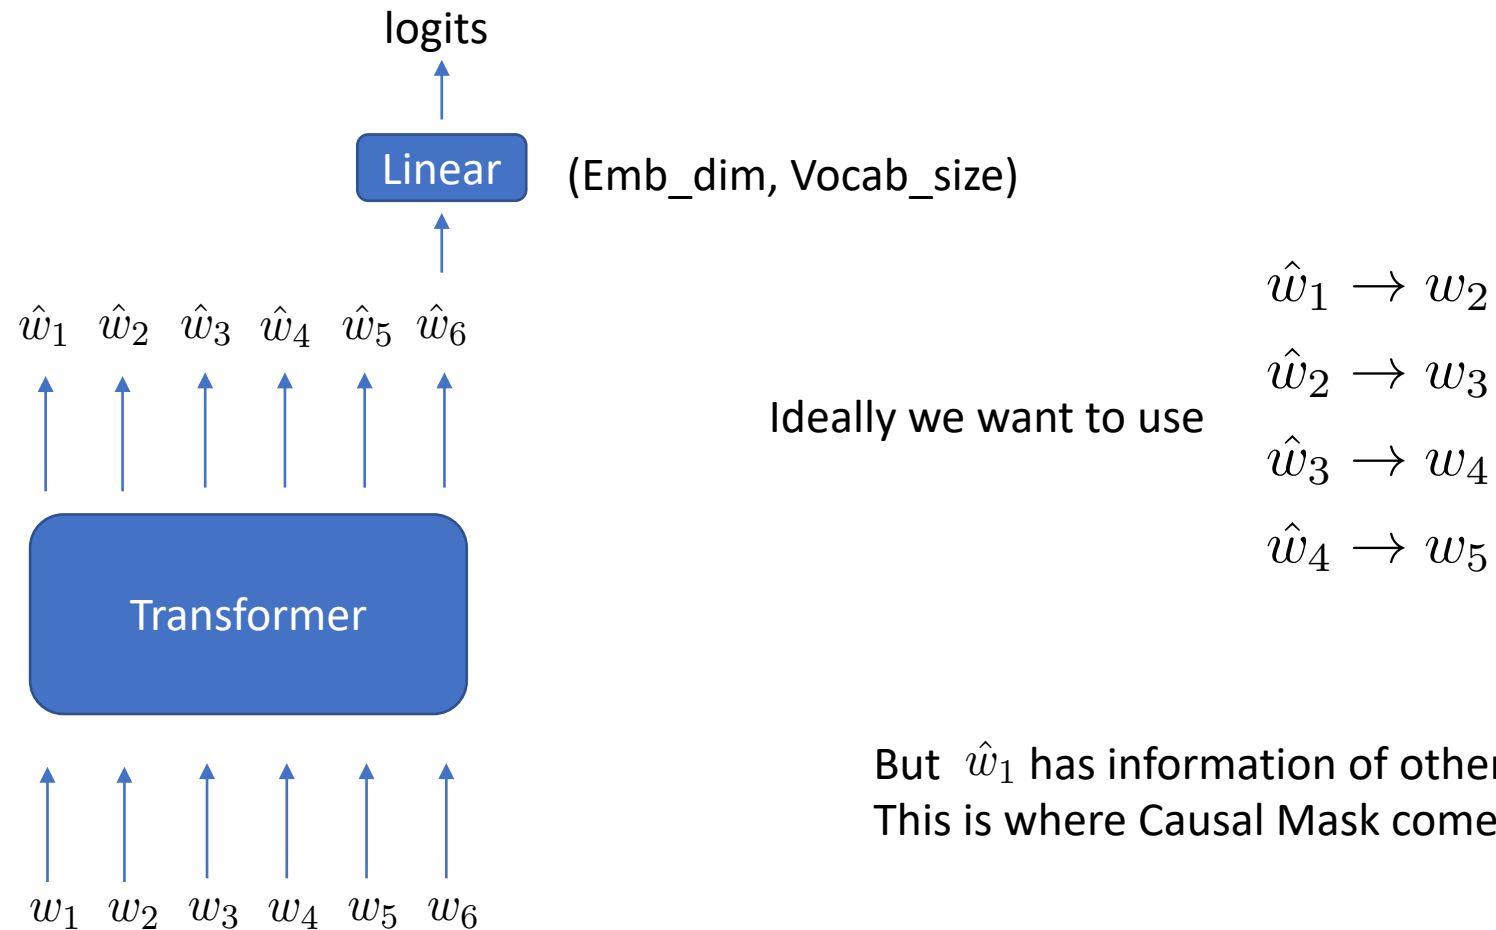
# How would we train a model to generate the next token/word?



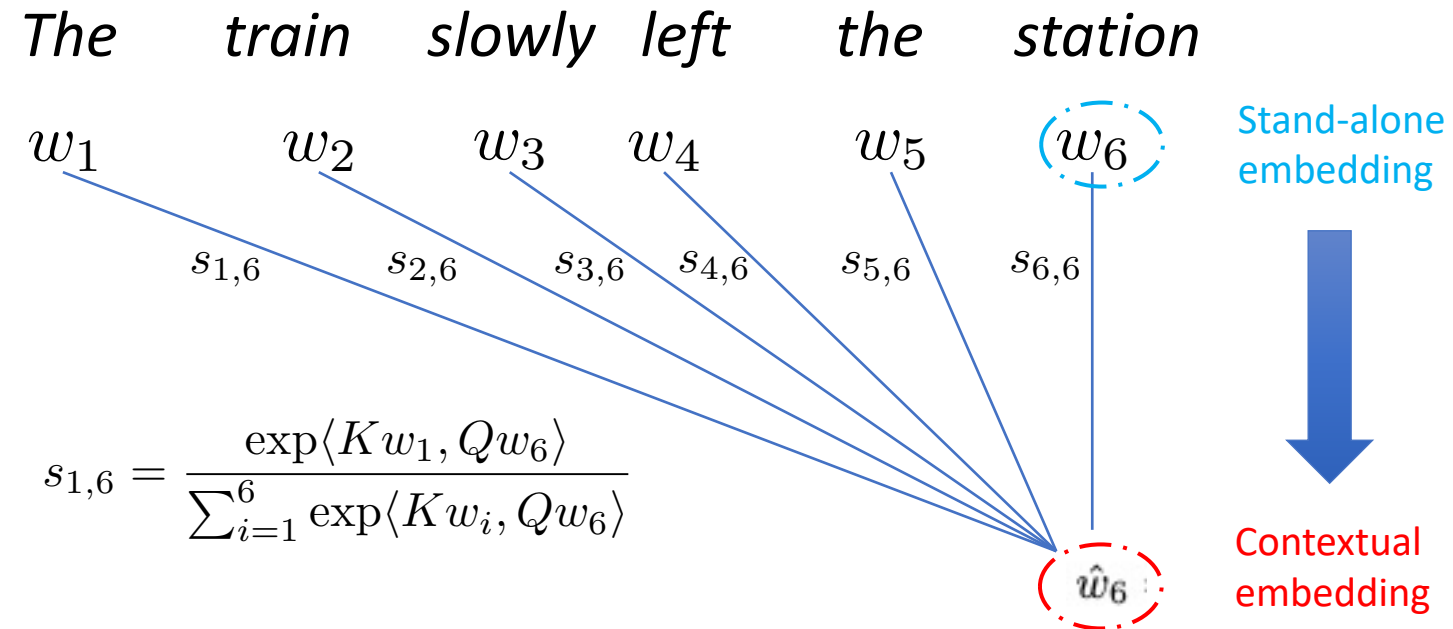
# How would we train a model to generate the next token/word?



# How would we train a model to generate the next token/word?



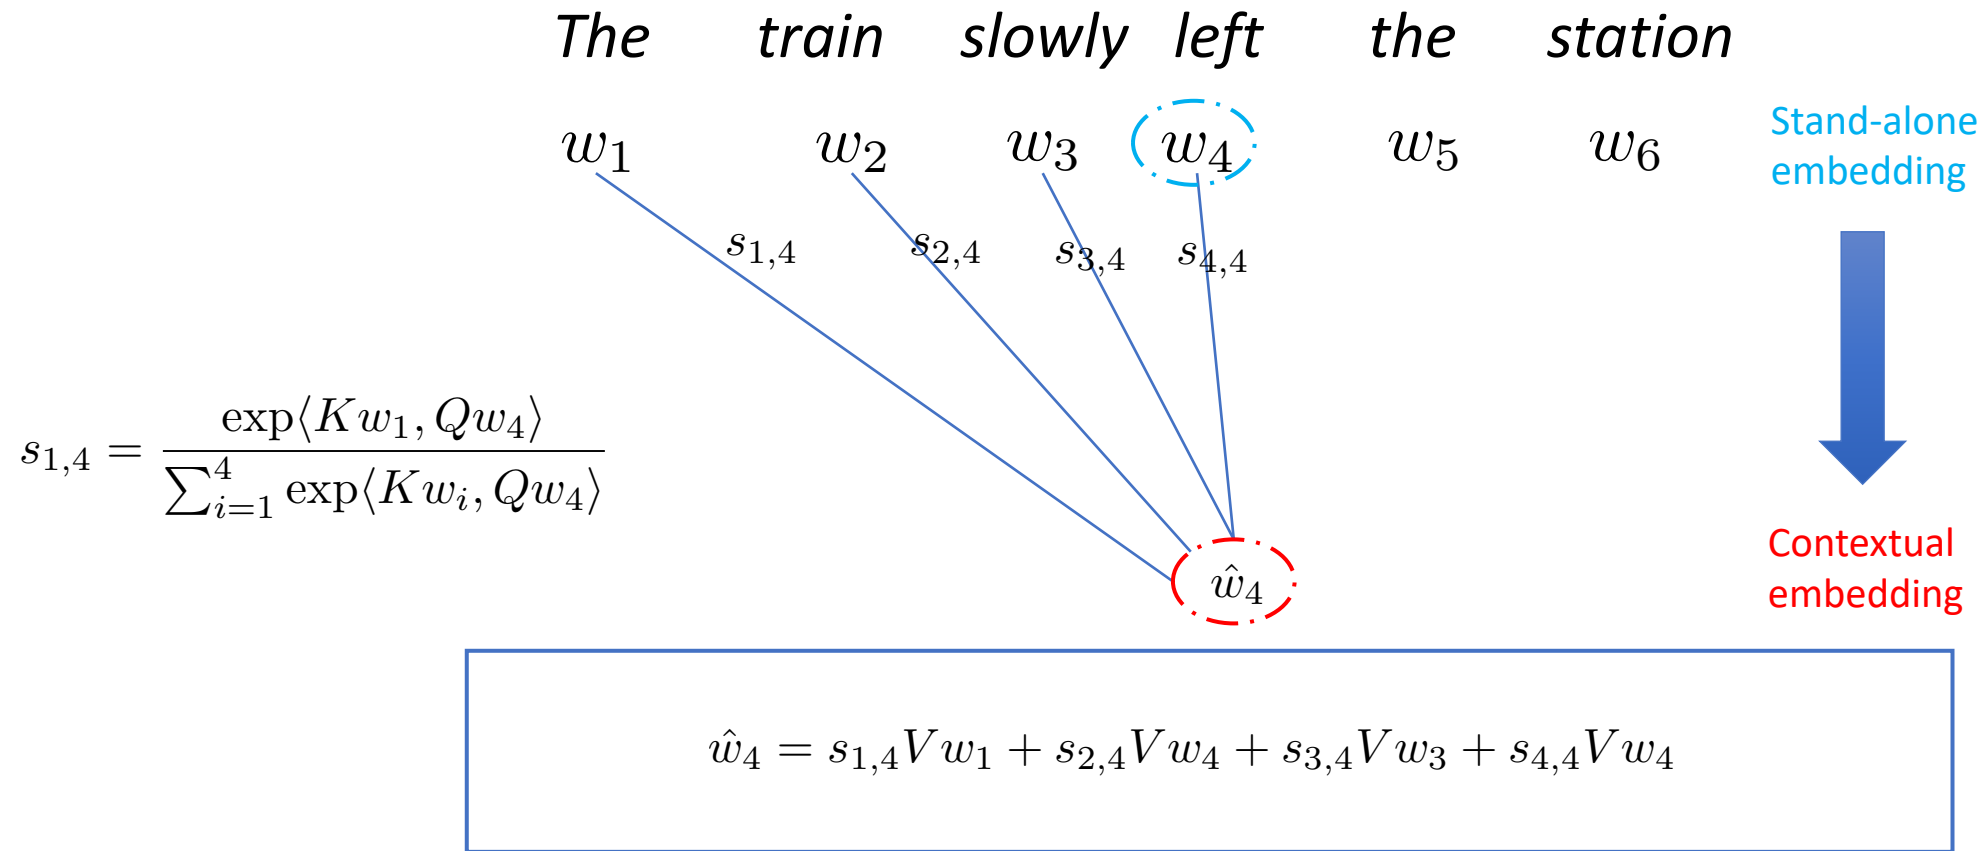
# In transformer, the information spread across tokens by Self-Attention



$$\hat{w}_6 = s_{1,6} V w_1 + s_{2,6} V w_2 + s_{3,6} V w_3 + s_{4,6} V w_4 + s_{5,6} V w_5 + s_{6,6} V w_6$$



# Causal Mask requires a token cannot use the information of a future token



# Causal Mask requires a token cannot use the information of a future token

$$A_{ij} = \text{mask}_{ij} \cdot \exp \langle K w_i, Q w_j \rangle$$

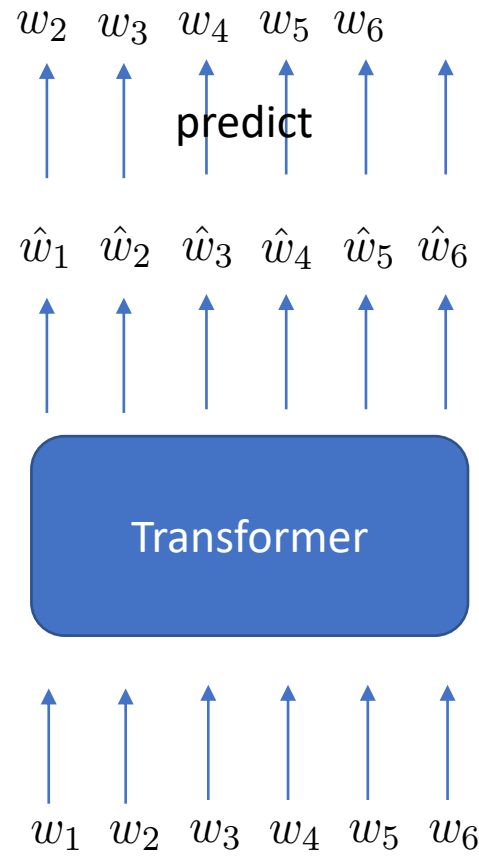
$$s_{i,k} = \frac{A_{ij}}{\sum_{i=1}^n A_{ij}}$$

	the	cat	sat	on	the
the	1				
cat	1	1			
sat	1	1	1		
on	1	1	1	1	
the	1	1	1	1	1

mask

Causal Self-Attention, Or Masked Self Attention

# With Causal Mask, *training* is much more efficient



All predictions/loss calculations  
can be done by one pass

# Autoregressive Language Model

With this training, essentially we obtain an oracle:  $p(x_k | x_{<k})$ , an autoregressive model for language

## This is what we called GPT!

GPT: Generative Pretrained Transformers

1. Transformer-based text-generation model
2. Pre-trained on **Massive** amount data

# Autoregressive Language Model

With this training, essentially we obtain an oracle:  $p(x_k | x_{<k})$ , an autoregressive model for language

GPT2 Model Architecture

Parameters	Layers	$d_{model}$
117M	12	768
345M	24	1024
762M	36	1280
1542M	48	1600

# Autoregressive Language Model

With this training, essentially we obtain an oracle:  $p(x_k | x_{<k})$ , an autoregressive model for language

GPT2 Model Architecture

Parameters	Layers	$d_{model}$
117M	12	768
345M	24	1024
762M	36	1280
1542M	48	1600

GPT2 Dataset

How would you collect high quality data from Internet?

# Autoregressive Language Model

With this training, essentially we obtain an oracle:  $p(x_k | x_{<k})$ , an autoregressive model for language

GPT2 Model Architecture

Parameters	Layers	$d_{model}$
117M	12	768
345M	24	1024
762M	36	1280
1542M	48	1600

GPT2 Dataset

Links from Reddit Posts that have at least 3 karmas

WebText: 45M links -> 8M docs (40GB)

# Autoregressive Language Model

With this training, essentially we obtain an oracle:  $p(x_k | x_{<k})$ , an autoregressive model for language

GPT2 Model Architecture

Parameters	Layers	$d_{model}$
117M	12	768
345M	24	1024
762M	36	1280
1542M	48	1600

GPT2 Dataset

Links from Reddit Posts that have at least 3 karmas

WebText: 45M links -> 8M docs (40GB)

Right Data + Right Model



# Recap: Transformer

- Transformer
  - Inputs: Input embedding + Position embedding
  - Transformer Layers
    - MultiHeadAttention
    - Regularizations
      - Dropout
      - LayerNorm
      - Residual Connection
    - Fast Forward Neural Networks
- Text Generation
  - Causal Mask for Transformer

# Assignment